

The Power to Choose the Counterfactual: Effective Demand, Intersubjectivity Collapse, and the Architecture of Global Deprivation

Flyxion
Independent Researcher

2026

Abstract

This essay argues that the coexistence of extravagant private consumption and severe national deprivation is not primarily a moral failing of individuals but the output of an epistemic architecture with a single underlying structure: counterfactual sovereignty, the unequal authority to determine both the feasible set of alternatives against which an action is judged and which member of that set serves as the comparator. Markets allocate resources according to effective demand rather than human need, and treat the former as sufficient authorization regardless of the latter. This same substitution reappears, in a different guise, whenever an intervention is judged successful merely because it produced something good rather than because it outperformed the reasonably available alternatives it displaced. Meritocratic self-narratives, philanthropic giving, infrastructure evaluation, and chokepoint capitalism all commit a version of this error, treating survivors, beneficiaries, or successfully cleared transactions as the whole of the relevant evidence while the excluded population, those who tried and failed, those who were not funded, those priced out of a chokepoint, disappears from view. Left unexamined for long enough, that excluded population disappears not only from institutional measurement but from moral perception itself: this is intersubjectivity collapse, a psychological endpoint reached only after an earlier and more fundamental measurement invisibility, in which unmet need first fails to register as effective demand at all. Using The Gambia and Cuba as empirically grounded cases whose divergent mechanisms, distributed debt-market risk management in one, deliberate geopolitical coercion entangled with domestic policy in the other, converge on a similar lived deprivation, the essay traces the recursive loop by which material separation produces institutional abstraction, institutional abstraction produces collapsed recognition, collapsed recognition produces normalized apathy, and normalized apathy reproduces the material separation that started the cycle. It closes by confronting, rather than resolving, the question of what personal obligation follows once deprivation is understood not merely as known but as occurring within institutions in which the observer is already a participant.

1 Introduction

Someone eating dinner on a terrace in a wealthy capital, planning a ten-thousand-dollar vacation to a resort with dependable electricity and imported food, is rarely thinking about a hospital in Banjul running on a diesel generator during a fuel shortage, or a family in Havana cooking over a wood fire because the grid has gone down for the third time in a week. This is not usually because the diner is cruel. It is because almost nothing in the architecture connecting the two of them is designed to make that connection perceptible, and because the economic system mediating both of their lives has already answered, silently and in advance, the question of whose claim on scarce global resources counts as real.

That silent answer is the starting point of this essay: effective demand is not equivalent to human need. A market can accurately, faithfully, without conspiracy or malfunction, report immense demand for luxury travel while simultaneously reporting negligible demand for rural electrification, because demand in the economic sense is not a measure of how much something matters to a human life. It is a measure of how much purchasing power stands behind a want. The vacation and the blackout are not unrelated facts that happen to coexist. They are two outputs of a single allocation procedure that was never designed to compare them in the first place.

The essay develops that claim across ten movements. It distinguishes effective demand from human need and traces what work that distinction does once followed to its consequences. It traces how material separation between persons becomes institutional abstraction, and how institutional abstraction produces what this essay calls intersubjectivity collapse, in a recursive loop rather than a one-directional causal chain. It identifies a shared epistemic error, the violation of what is here called the counterfactual discipline, running through meritocratic self-narrative, philanthropic giving, and infrastructure evaluation, extends that error to the prior question of which capacities institutions allow to become visible at all, and argues that chokepoint capitalism is best understood as an instance of the same error operating at civilizational scale, with genuine antecedents in colonial commodity and currency systems. It grounds the argument in two cases, The Gambia and Cuba, using debt, electrification, food-import, and historical data rather than illustration alone, and takes seriously the objection that domestic governance, not external architecture alone, shapes both outcomes. It closes by returning to the question the essay cannot honestly avoid: what a person who benefits from this architecture, who can in fact afford the vacation, owes to a world in which the same architecture produces the blackout.

2 Effective Demand Is Not Equivalent to Human Need

Markets are frequently described, by both defenders and critics, as neutral aggregators of preference. Whatever people want most, weighted by how much they are willing to pay, is what gets produced. The description is accurate as far as it goes. Its danger lies in what it leaves silent: willingness to pay is not a measure of how much a person needs something, only of how much purchasing power they can direct toward it. A wealthy individual's preference for a third vacation home registers in markets with a strength that a rural clinic's need for a reliable power supply, absent some external subsidy, typically cannot approach, not because the vacation home matters more to human welfare, but because the two claims enter the market through radically different quantities of backing money.

This is not a novel economic observation; the distinction between effective demand and need has a long history in welfare economics, running through debates over utilitarian aggregation and the diminishing marginal utility of income that occupied economists such as Pigou and, later, critics of Pareto-efficiency as a normative standard. What this essay adds is not the observation itself but an insistence on tracing its consequence all the way through: if markets systematically privilege demand backed by money over need unbacked by money, then a global economy dominated by market allocation will systematically privilege the discretionary consumption of the wealthy over the survival needs of the poor, and it will do this continuously, without any single decision-maker ever choosing that outcome. Nothing paradoxical or conspiratorial has occurred at the level of any individual transaction. The paradox appears only when the aggregate output is compared against what human need alone would recommend, a comparison the market itself has no mechanism for making, because it was never asked to make it.

This distinction reframes what intersubjectivity collapse actually consists of. It is tempting to treat collapse as a purely psychological phenomenon, a numbing of empathy under the weight of scale, examined in the following section. The deeper form of collapse is institutional before it is psychological: it occurs whenever a person's subjectivity enters a decision procedure as purchasing power, and another person's enters as an externality, a risk category, an aid statistic, or nothing at all. The excluded person has not vanished physically. Their claim upon the world's finite stock of labor, energy, land, and material has vanished from the representation that the decision procedure actually consults. A price system is, among other things, a machine for converting persons into demand curves, and demand curves, unlike persons, have no built-in mechanism for registering desperation that lacks money behind it.

The full sequence runs from measurement to experience rather than the reverse: a need first fails to register as effective demand, which produces institutional invisibility in the

specific systems, credit models, aid formulas, trade statistics, tasked with representing that need; institutional invisibility then produces the experiential distance described in the next section, since an observer's information environment is itself assembled largely from what those systems chose to represent; and experiential distance finally produces reduced political salience, closing the sequence by removing the pressure that might otherwise have corrected the original measurement failure. Psychological invisibility, in other words, follows an earlier and more fundamental measurement invisibility. Treating apathy as the origin of the problem, rather than as the last link in a chain that begins in section two's distinction between demand and need, mistakes the symptom for the cause.

3 Material Separation, Institutional Abstraction, and the Collapse of Intersubjectivity

It is worth resisting a simpler, more familiar version of this argument, in which structural deprivation is largely a downstream consequence of psychological indifference: people would act differently if only they felt enough, and apathy is the missing variable holding back an otherwise available solution. The relationship this essay proposes runs at least as strongly in the opposite direction, and forms a loop rather than a line.

Material separation comes first, in the sense of being the precondition rather than the consequence. Long before any psychological process occurs, the tourist and the family rationing cooking oil are already separated by geography, currency, legal jurisdiction, and by the physical infrastructure of ports, cables, and bank clearing systems that route almost none of their respective daily lives through contact with each other.

Institutional abstraction follows directly from material separation, not as an optional add-on but as an operational necessity for any institution that must act at the resulting scale. A correspondent bank cannot process a transaction as an encounter between two named individuals; it must convert the transaction into a risk score, a jurisdiction code, a compliance flag. A credit rating agency cannot evaluate a sovereign borrower as a nation of specific families; it must convert the nation into a small number of macroeconomic ratios. Abstraction of this kind is not a moral failing on the part of any individual compliance officer or analyst. It is what the institutional role requires of anyone who occupies it, and it would be required of a far more compassionate occupant of the same role just as strongly.

Intersubjectivity collapse is what happens to observers, including the officials and analysts themselves, once the representation they work with has been thoroughly purged of anything that would register the other party as a subject rather than a score. The phenomenological tradition running from Husserl through Levinas grounded ethical obligation in the encounter with a face, a presence that resists reduction to any category. Bracketing the

metaphysical commitments of that tradition, its empirical descendant in social psychology is directly relevant here. Paul Slovic's research on what he termed psychic numbing found that willingness to help and reported emotional concern fail to scale, and in some conditions actually decline, as the number of people at risk grows from one identified victim toward an anonymous multitude, a pattern he called the collapse of compassion. The sociologist Luc Boltanski's analysis of mediated suffering adds an institutional layer to the same finding: modern spectators of distant suffering, encountering it primarily through news media, are placed in what he called the position of the spectator, structurally distant from both the sufferer and any means of direct action, a position that itself shapes what response, if any, becomes psychologically available. Slovic and Boltanski together describe the human cognitive and social endpoint of a process that begins institutionally, in the abstraction described above, not the whole of its cause.

Normalized apathy is the stabilized affective state that results once collapse has been sustained long enough to become background rather than event. It is a specific and produced response, not a null result, and it should not be treated as equivalent to cruelty; a person who feels nothing about a supply chain's distant human cost is not, typically, someone who has decided that cost does not matter. They are someone for whom the cost was never rendered, cognitively, as a cost borne by a subject.

Political reproduction closes the loop. Normalized apathy removes the public pressure that might otherwise force reform of the institutions generating material separation in the first place, correspondent banking rules, sanctions carve-outs, sovereign debt restructuring norms, allowing the architecture producing that separation to persist and, in many cases, deepen. The deepened separation then generates a fresh round of institutional abstraction, a fresh round of collapse, and a fresh round of apathy.

The loop matters because it changes where intervention is likely to be effective. Efforts aimed solely at increasing empathy in distant publics, more vivid imagery, more urgent framing, are fighting an architecture that requires abstraction as an operational prerequisite at its point of control; a newly sensitized public has no natural point of contact with the compliance officer or ratings analyst whose daily abstraction reproduces the chokepoint regardless of public sentiment. Effective intervention is more likely to require altering the operational logic of the institutions themselves: transparency requirements on correspondent-banking de-risking decisions, human rights impact assessments built into sovereign credit methodologies, and humanitarian exemptions in sanctions regimes robust enough to survive banks' risk aversion rather than existing only on paper.

4 The Counterfactual Discipline

4.1 A Shared Epistemic Error

Consider four apparently unrelated situations: a public figure attributing his success to a disciplined morning routine, a foundation announcing that its program helped fifty thousand people, a city celebrating a thirty percent drop in bicycle collisions after a new bike lane, and an economist observing that a chokepoint institution successfully cleared ninety-nine percent of the transactions submitted to it. Each of these statements shares the same defect: it reports an observed outcome without specifying, or in some cases without even acknowledging the existence of, the relevant counterfactual against which that outcome should be judged.

Empirical science exists partly to prevent exactly this defect. A researcher does not conclude that an intervention caused an effect merely by observing the effect after the intervention. The standard procedure requires selecting features to measure, justifying why those features are good proxies for the underlying phenomenon of interest, specifying a baseline or null model describing what the data would look like in the intervention's absence, and setting a threshold of significance before treating any observed difference as meaningful rather than as noise. The randomized controlled trials pioneered in development economics by Abhijit Banerjee, Esther Duflo, and Michael Kremer, work recognized by the 2019 Nobel Memorial Prize in Economic Sciences, exist precisely because informal program evaluation, of the kind foundations and governments had relied on for decades, routinely failed this discipline: interventions were judged successful by comparing recipients to no baseline at all, or to a baseline chosen after the fact to flatter the result. The demand that an intervention be compared against a randomly assigned control group, rather than against an arbitrary or self-selected comparison, is simply the formal version of the discipline this essay is arguing has been abandoned everywhere outside a narrow band of specialist practice.

Social life routinely dispenses with this discipline, and it does so in a way that is not randomly distributed. It is disproportionately the powerful, those with the resources to broadcast a narrative and the least incentive to interrogate it, who get to decide which counterfactual their success or their generosity will be judged against. The point sharpens further once the comparator is recognized as a choice among many rather than a single control condition: for any action a there is typically a feasible set of alternatives $\mathcal{C}_F(a)$ that were reasonably available at the time of decision, not a lone alternative, and the comparator actually used in evaluation should be selected by a rule declared independently of how a turns out, using only information available before the outcome was known. An actor permitted to select, after the fact, the weakest member of that set makes almost any action look like an improvement; an actor permitted to select, after the fact, whichever member of

that set happened to perform best commits the mirror-image error, since one alternative can succeed spectacularly through sheer unpredictable circumstance and thereby become an unreasonable standard against which everything else, including *a*, is then judged a failure. The general principle is therefore not that an action must be compared against the worst case, nor that it must be compared against the best case, but that the comparator must be justified independently of the observed result. Appendix B formalizes this distinction. The deeper claim running through the remainder of this section, then, is not merely that power lets an actor choose what happens, but that it lets an actor choose what the comparison class for what happened will be, a capacity this essay calls counterfactual sovereignty: unequal authority over the space of alternatives, and over the rule for selecting among them, by which actions and institutions are subsequently declared necessary, successful, generous, efficient, or just. Effective demand, meritocracy, the talent reserve, philanthropy, infrastructure evaluation, and chokepoint capitalism, treated separately in the sections that follow, are not six analogous problems so much as six instances of this single operation.

4.2 Meritocracy as Survivorship Bias

Michael Sandel's *The Tyranny of Merit* argues that meritocratic societies encourage their winners to regard success as deserved and their losers as responsible for failure, with corrosive effects on social solidarity. That argument is moral and political. There is a methodological argument sitting underneath it that deserves separate statement, because it does not depend on any claim about what winners deserve, only on a claim about what can validly be inferred from their example.

Suppose several million people adopt some large variety of morning routines, risk postures, educational strategies, and networking habits, including a considerable amount of sheer trial and error. In a sufficiently large population, some very small subset will achieve spectacular success purely as a matter of chance, in the same sense that flipping millions of coins guarantees a run of ten consecutive heads somewhere in the data even though no coin is biased. Looking backward from that successful subset and asking what distinguishes them from everyone else will nearly always turn up some apparently distinctive trait, because with enough traits examined after the fact, some will correlate with the outcome by chance alone. What that retrospective search cannot establish, without further information the search itself does not provide, is whether the trait caused the success, since establishing causation requires knowing the trait's prevalence among the much larger population that possessed it and did not succeed, its prevalence before success occurred, and what a reasonable model of chance variation would have predicted in its absence.

A public figure who explains success by pointing to relentless work, or to a specific moral outlook, the kind of explanation regularly offered by prominent entrepreneurs discussing

personal habits or by founders framing their success in the language of a personal mission, has not thereby established that explanation merely by accurately reporting that he worked relentlessly, or held that outlook, and also succeeded. Millions of people who worked equally hard, or held similar convictions, did not succeed, and any honest accounting of causation would need their outcomes as the denominator. The selection here operates twice, not once: first over outcomes, in that a winner is selected from a much larger population by a process substantially driven by chance and opportunity, and then over explanations, in that the winner, being the one still present and narrating, possesses disproportionate authority to determine which of his traits become the socially legible story of why he won, while the much larger comparison population of equally diligent people who did not receive the same sequence of opportunities remains silent and uncounted. Survivorship bias becomes a matter of public consequence, not merely a technical statistical footnote, at the moment survivors acquire the cultural authority to explain the very selection process that produced them, an authority meritocratic culture grants nearly automatically to anyone who has won.

4.3 Scarcity, Sprezzatura, and the Talent Reserve

The selection problem described in the preceding subsection has a counterpart on the input side, before any winner has emerged for meritocratic narrative to explain. Institutions do not merely misattribute the causes of success once success has occurred; they also, and prior to that, systematically underestimate how much capacity exists to succeed in the first place, because what they can measure is only ever revealed capacity, not available capacity, and available capacity is in turn only a fraction of what might be called possible capacity. Writing T_{visible} for the talent an institution's credentials, employment records, and formal demonstrations actually register, $T_{\text{available}}$ for the capacity that exists in the population and could be exercised under different conditions, and T_{possible} for the full capacity a population could in principle develop given sufficient opportunity, institutions routinely reason as though $T_{\text{visible}} \approx T_{\text{available}} \approx T_{\text{possible}}$, treating an observed shortfall $S_{\text{observed}} \ll S_{\text{potential}}$ as though it described the population's actual ceiling rather than the ceiling of what current measurement procedures happen to detect. Claims such as a shortage of qualified engineers, an absence of local expertise, or a lack of technical talent in a given country can become self-confirming in exactly this way, since the institutions responsible for detecting capability recognize only particular credentials, employment histories, languages, or geographic positions, and treat everything outside that narrow aperture as though it did not exist.

This gap between visible and available capacity is not always imposed from outside; it is sometimes actively maintained by the people who possess the capacity, for reasons that are entirely rational given the incentives they face. Baldassare Castiglione's sixteenth-century

account of *sprezzatura*, the cultivated appearance of effortlessness that conceals the labor, calculation, and practice behind a skill, describes a stylistic version of this concealment. A more consequential version operates beneath the level of style: a person who demonstrates competence at level q while privately capable of $q + r$ is not necessarily failing to reveal r through incapacity or modesty. If revealing r reliably produces additional uncompensated demands, expectations, exposure to competition, or loss of bargaining power, without a proportional increase in autonomy, compensation, or security, then withholding r is a rational protection of a reserve of agency rather than a failure of disclosure. Revealed capacity, under these conditions, systematically understates available capacity, and an institution that reads revealed capacity as though it exhausted available capacity will conclude that a shortage exists precisely where a reserve is being deliberately protected.

Call the capacity that exists but has not become institutionally legible the talent reserve, the structural inverse of brain drain. Brain drain describes a system losing capacity that has already been detected, credentialed, and recognized, typically to wherever that recognition is better rewarded. Talent reserve describes capacity that never became institutionally legible to begin with, not because it does not exist but because the credentials, employment histories, market positions, or geographic locations through which an institution recognizes talent never had occasion, or incentive, to register it. The distinction between the two states, and a third describing what happens when latent capacity becomes exercisable, can be stated compactly: brain drain is the departure of already-visible capacity; talent reserve is capacity that remains present but stays latent; talent activation is the process by which latent capacity becomes exercisable without necessarily requiring departure at all. Appendix D states this distinction as a measurement equation, with visible capability as a function of latent capability and institutional opportunity, and the reserve as the population-wide sum of what that function fails to register.

The connection to the counterfactual discipline developed earlier in this section is direct. Suppose a country selects its ten highest-scoring students for scholarships, and all ten subsequently become successful engineers. The weak counterfactual asks whether these ten would have become engineers without the scholarship, and against that comparison the program looks like an unambiguous success. The stronger and more informative counterfactual asks what would have happened had the same resources instead been used to create conditions under which ten thousand people could reveal capacities that current institutions cannot currently detect. The first comparison evaluates whether selecting already-legible talent was effective. The second tests whether selection was the right architecture to begin with, as opposed to activation, and it is precisely this second comparison that a scarcity narrative, by treating T_{visible} as though it were T_{possible} , prevents from ever being posed.

This connects to the case studies in a way that deserves stating plainly rather than left

as an aside. Describing a country as resource-poor, skills-constrained, or lacking human capital can silently convert an observation about currently realized capacity into an assertion about a population's intrinsic capacity, and once that conversion has occurred, the resulting deprivation becomes evidence for its own explanation: constrained infrastructure produces fewer legible opportunities to demonstrate skill; fewer legible opportunities produce fewer conventionally credentialed experts; their absence is read as a skills shortage; and the shortage is then invoked to justify continued concentration of scarce investment in the small number of people who did manage to become legible, rather than in the conditions that would let a much larger population reveal what it already contains. The loop closes on itself: material scarcity restricts opportunity, restricted opportunity lowers the rate of capability revelation, low revelation is read as an intrinsic talent scarcity, perceived scarcity justifies concentrating investment in an already-selected few, and that concentration reproduces the material scarcity it was meant to explain. Institutions operating under this loop do not merely choose a counterfactual after an outcome has already occurred, as in the meritocracy and philanthropy cases above. They partly determine, through what their measurement procedures are built to detect, which unrealized counterfactual capacities are ever allowed to become observable in the first place, which is counterfactual sovereignty operating one step further upstream than Section 4.1 initially described it.

4.4 Philanthropy as Selection and Diversion

The conventional critique of large-scale philanthropy concerns influence: a small number of wealthy donors acquire outsized power to shape public priorities that were once, at least nominally, subject to democratic deliberation. This essay's counterfactual argument supports that critique but adds a distinct point about how philanthropic success is measured.

A funded intervention is almost always evaluated against the baseline of its beneficiaries receiving nothing at all. Against that baseline, nearly any competently administered program looks like a success, because doing something is compared only to doing nothing, the weakest and least informative comparison available. But allocating a large sum to cause A is simultaneously a decision not to allocate it to causes B, C, or D, not to fund public provision through taxation, not to support structural reforms that might reduce the production of the deprivation the donation addresses, and not to alter the chokepoint architecture that generated the fortune being donated in the first place. Giving is inseparable from diverting; every unit of resource directed toward one trajectory is a unit withheld from every other trajectory that same resource could have supported.

The appropriate evaluative question is therefore not whether a philanthropic program helped, a criterion so permissive that it is nearly always satisfied, but relative to which feasible alternatives, measured against which justified variables, and against which baseline

the program performed better. Without that specification, an apparent philanthropic success may be closer to a false positive generated by an artificially weak null hypothesis than to evidence that the resources were well allocated. This is not an argument against giving; it is an argument that giving's moral accounting has been allowed to grade itself against the easiest available exam.

The point is stronger, in fact, than the question famously posed by Peter Singer, of what an affluent individual ought to do given the existence of preventable suffering elsewhere. Singer's question already accepts that the affluent individual is the one who gets to decide, presenting the choice as roughly a vacation against a donation. The antecedent and more uncomfortable question this essay is pressing is why any single actor, however generous, should possess so much authority over the admissible alternatives in the first place. The realistic counterfactual set for a large philanthropic sum is not {donate, do not donate}; it includes taxation, universal public provision, redistribution, alternative ownership arrangements over the assets that produced the surplus, debt relief, reformed trade institutions, direct unconditional transfers, and arrangements under which neither the donor nor any other private actor retains discretion over who receives the surplus at all. A philanthropic program can therefore genuinely improve the realized world relative to no intervention while simultaneously legitimating an institutional arrangement that keeps better members of that counterfactual set from ever being seriously considered. These two claims are not in tension; recognizing that they can both be true at once is precisely what evaluating philanthropy only against a do-nothing baseline prevents.

4.5 Metrics, Baselines, and the Bike Lane

The same error recurs at a much smaller and more concrete scale, which is useful precisely because it keeps the argument from remaining abstract. Consider a deliberately hypothetical case, chosen for illustration rather than as a claim about any actual city or study: a city installs a protected bike lane and subsequently reports a thirty percent reduction in bicycle collisions. If automobile collisions rise by ten percent, emergency response times lengthen because a traffic lane was removed, and displaced traffic increases pedestrian conflict at nearby intersections, then celebrating the bicycle-collision figure alone is a metric-bound false positive: true on its own terms, and misleading about the intervention's actual effect on total road safety, because the observer selected, in advance or in retrospect, the one metric on which the intervention happened to succeed. The relevant objective was never a single group's collision count but total expected harm across every road-user group,

$$H = \sum_{g \in G} \sum_{k \in K} P_{gk} L_{gk},$$

where g ranges over road-user groups, cyclists, drivers, pedestrians, emergency responders, k ranges over incident types, and P_{gk} and L_{gk} are the probability and severity of incident type k for group g . An intervention is an improvement only if $H_{\text{after}} < H_J$ for a justified comparator H_J in the sense of Appendix B, not merely if the single term corresponding to bicycle collisions falls. Reporting a decline in one P_{gk} while H overall may have risen is exactly the selective-metric error this section has been describing throughout.

A second and subtler problem concerns who is positioned to design the intervention in the first place. A street contains its painted lines and posted signs, but it also contains an accumulated, largely unwritten protocol: where drivers actually merge, where pedestrians and cyclists have learned to expect each other, what a nearly worn-away lane marking used to signify to people who have used that street for years. This situated knowledge is distributed across long-term users rather than represented in any engineering drawing, and it frequently survives even after the formal markings that once encoded it have faded, which is why longtime residents can navigate a street with worn paint while visitors cannot. A planner equipped with an extremely detailed geometric model of a street can nonetheless lack the behavioral knowledge that its regular users possess collectively, and a planning process that treats the engineering drawing as the whole of the relevant reality, while never consulting people who have actually driven and cycled that specific street over years, is vulnerable to producing formally accurate designs that are behaviorally false. This is a smaller-scale instance of the abstraction problem described in section three: a representation can be technically precise and still have discarded the very information that made it useful.

5 Chokepoint Capitalism: From Colonial Commodity Systems to Modern Bottlenecks

The term chokepoint capitalism, developed by Rebecca Giblin and Cory Doctorow, originally described intermediaries in cultural and digital markets, publishing platforms, streaming services, app stores, that position themselves at unavoidable bottlenecks between creators and audiences and extract disproportionate value from both sides regardless of how competitive the surrounding market appears. The same structural logic, control of a narrow passage that cannot practically be avoided, operates at the scale of nations, and it did not appear from nowhere.

Colonial monopoly trading companies controlled ports and shipping routes in much the way container shipping alliances now control global freight capacity. Colonial currency boards, which pegged colonial currencies to the metropole's and required reserves to be held in metropolitan banks, extracted seigniorage and control simultaneously, anticipating the modern dependency of global trade on dollar clearing and correspondent banking.

Colonial administrations rated and ranked territories for creditworthiness using language about institutional immaturity strikingly similar to that used by modern sovereign credit agencies to justify unfavorable terms. When the International Monetary Fund and World Bank extended structural adjustment programs to newly independent states across the 1980s and 1990s, conditioning further lending on currency devaluation, subsidy removal, and public-sector retrenchment, the underlying logic, external control over the terms on which a formally sovereign state could access capital, was continuous with the colonial currency-board arrangements it formally replaced, even as the institutional vocabulary shifted from empire to development finance. The technical vocabulary has changed across each of these transitions; the structural position, controlling a passage that a formally sovereign but practically dependent territory cannot avoid using, has not.

Four contemporary chokepoints matter most for the cases that follow. Reserve currency and payment rails: the overwhelming majority of international trade is invoiced and settled in dollars and clears through a small number of correspondent banks and the SWIFT messaging system, an unavoidable passage through which sanctions, transaction fees, and unilateral bank de-risking decisions can constrain a country's imports without a single ship being physically blockaded. Shipping and logistics: a small number of alliances and a handful of critical maritime corridors carry the overwhelming share of global trade, and nations without deepwater ports or the diplomatic weight to negotiate favorable terms pay a persistent premium on every imported good. Debt and credit rating: a small number of rating agencies treat developing-nation debt as a distinct, higher-risk asset class, and a downgrade raises borrowing costs, which worsens the fiscal position that triggered the downgrade, a feedback loop that constrains infrastructure financing long after any initiating shock has passed. Sanctions and extraterritorial enforcement: because so much of global finance routes through institutions exposed to United States regulation even when neither transacting party is American, sanctions can be enforced far beyond their nominal jurisdiction, with third-country banks often declining to process a technically legal transaction simply because the compliance risk of even appearing to test the boundary is judged not worth the business.

None of these four domains requires anyone to intend deprivation as an outcome, which is precisely what makes the resulting apathy described in section three so stable. Correspondent banks manage regulatory exposure, rating agencies price risk, shipping alliances optimize routes, each acting in a manner that is individually defensible and rarely traceable to a single identifiable decision-maker, and the combined effect is a durable architecture in which certain nations absorb systemic costs that appear, from any single actor's vantage point, as no one's deliberate choice.

6 Two Case Studies

6.1 The Gambia

The Gambia's modern economic structure still bears the imprint of a colonial-era specialization in groundnut export, imposed under British administration and inherited largely intact at independence in 1965. That monoculture left the country dependent on a single volatile commodity rather than a diversified domestic industrial base capable of generating independent foreign exchange, and it set a template that subsequent decades of development finance have reproduced in a different form: dependence on external actors, first a single export buyer, later external creditors and donors, for the country's ability to finance its own infrastructure. An Economic Recovery Programme launched in 1985 under IMF and World Bank guidance imposed currency devaluation and public-sector austerity in exchange for continued lending, a pattern that recurred in various forms through the decades that followed, including under the government that took power after the 1994 coup and the democratic transition of 2017.

The current figures illustrate the constraint concretely rather than abstractly. The Gambia's total government debt stood at roughly seventy-two to eighty percent of GDP through 2023 and 2024, with external public debt alone around forty-six to fifty-three percent of GDP across the preceding decade, and debt service consuming close to a quarter of the national budget [11, 9]. Tourism receipts and remittances from Gambians working abroad each account for roughly a fifth of GDP [9], meaning the country's two largest sources of foreign exchange are, respectively, a sector highly exposed to European economic cycles it does not control and a flow dependent on the wages of citizens working outside the country's own economy. Access to electricity nationally reached roughly sixty-seven percent of the population by 2023 [12], meaning close to a third of the country still lacks it, and the national electricity utility's persistent financial strain has continued to weigh on the government's fiscal balance even as generation capacity has grown [10].

Tourism's role illustrates the luxury-and-deprivation coupling with unusual directness. Coastal resorts serving European visitors typically maintain reliable power, often through dedicated generator capacity that exceeds what surrounding communities draw from the national grid, because tourism revenue is the least chokepoint-constrained source of foreign exchange available to a debt-burdened state, and investment therefore flows there disproportionately. This is not a conspiracy against the Gambian countryside. It is the predictable output of a state rationally directing scarce reliability toward the sector best positioned to service its dollar-denominated debt, given the debt-service obligations described above, a rationality that itself depends entirely on the chokepoint architecture that made debt service the state's binding constraint in the first place.

6.2 Cuba

Cuba presents a causally entangled rather than a singular story, and the essay is stronger for refusing to flatten it into one. The United States embargo, in place in some form since 1960 and substantially tightened by the Cuban Democracy Act of 1992 and the Helms-Burton Act of 1996, the latter of which extended enforcement extraterritorially by allowing lawsuits against foreign companies deemed to be trafficking in property expropriated from American claimants, restricts direct trade and, through mechanisms resembling the sanctions dynamics described in section five, discourages third-country banks and firms from transacting with Cuba at all for fear of exposure to the United States financial system. A period of partial normalization under the Obama administration between 2014 and 2016 was substantially reversed after 2017, and further tightened again in the mid-2020s. But the embargo operates alongside, and interacts with, a set of domestically rooted vulnerabilities: a centrally planned economy that historically limited the foreign direct investment that could have diversified the country's energy and food supply, a Special Period of severe economic contraction following the 1991 dissolution of Soviet subsidies, decades of subsequent dependence on discounted Venezuelan oil, and a currency structure in which the government must purchase almost all imports using scarce foreign-currency reserves rather than the domestic peso.

The scale of the resulting strain is severe by any measure. Cuba experienced repeated nationwide and near-nationwide blackouts across 2024, 2025, and into 2026, driven by fuel shortages, a lack of spare parts for aging plants, and outright failures such as the October 2024 collapse of the Antonio Guiteras plant that triggered a full national blackout [13, 14]. Venezuelan oil shipments, which had covered on the order of a quarter of Cuba's daily oil deficit through 2025 [17], faced the prospect of falling toward zero following the January 2026 capture of Venezuelan leader Nicolás Maduro, a development Cuban officials and independent analysts alike warned would deepen an already severe fuel crisis [18]. Cuba imports an estimated seventy to eighty percent of its domestic food requirements, with most of that volume directed toward the social protection programs that keep the most vulnerable fed [15], and a 2024 United States Department of Agriculture assessment estimated that more than a tenth of the population fell short of minimum daily caloric intake that year, a figure that rose considerably higher under adjusted GDP assumptions [16]. Both the Cuban government and independent analysts agree that Venezuelan dependency and limited domestic energy diversification are genuine contributing causes, even as they disagree sharply about the relative weight of the embargo itself.

What distinguishes Cuba from The Gambia is not the severity of the outcome but the nature of the chokepoint. Gambia's debt architecture is largely the emergent output of ordinary risk-averse behavior by lenders, raters, and donors, none of whom individually designed a system to constrain Gambia specifically. Cuba's embargo is a chokepoint con-

structured and maintained as explicit policy, with pressure on the Cuban economy as one of its stated objectives, even though its effects are compounded by domestic decisions the Cuban government itself has made and continues to make. Whatever conclusion one reaches about the embargo's justification, a question this essay does not attempt to settle, it cannot be described as an unintended byproduct in the way that Gambia's debt structure largely can be.

6.3 What the Comparison Shows

The two cases test the theory rather than merely illustrating it, because they diverge in exactly the dimension that matters: intentionality of the chokepoint. Gambia shows that deprivation does not require anyone to want it; an architecture of ordinary risk management, applied to a small commodity-dependent economy, can reproduce scarcity indefinitely without a villain. Cuba shows that even where a chokepoint is deliberately constructed, its effects remain entangled with domestic choices in a way that resists any monocausal account, and a serious version of this essay's argument has to hold both of those truths simultaneously rather than collapsing toward whichever one is more rhetorically convenient. The pairing is instructive precisely because the mechanisms diverge so sharply, distributed incentive structures and ordinary debt-market risk management in one case, deliberate geopolitical coercion entangled with domestic energy policy in the other, while the phenomenology converges: blackouts, food insecurity, and medicine shortages that a tourist on a Gambian beach or transiting a Havana airport lounge has little structural reason ever to connect to the architecture financing their own trip. If both cases arrived at deprivation only through identical, clearly intentional coercion, the architectural claim of this essay would be far less interesting, since it would reduce to an argument about the specific actors responsible for that one coercive policy. That divergent mechanisms recurrently produce a convergent lived outcome is what supports the stronger claim that a system need not intend a given outcome in order to reliably reproduce it.

7 Objections and Complexities

The most serious objection to this essay's argument is that it risks understating domestic governance as a cause of both countries' strain, a risk this essay has tried to resist by presenting figures and causal claims that explicitly include debt-service ratios shaped by domestic borrowing decisions, and Cuba's own energy-diversification failures, rather than treating external architecture as the sole variable. The claim being made is not that governance is irrelevant. It is that governance failures are chronically overweighted in public discourse relative to chokepoint architecture, not because governance failures are smaller, but because

they are more visible: they attach to a named leader, they appear in a specific news cycle, they have the dramatic shape a story requires. A credit downgrade triggered by a rating agency's risk model, or a bank's routine de-risking decision, rarely generates an equivalent story, even when its fiscal effect on hospital funding is larger than any single act of domestic corruption.

A second objection concerns agency: does emphasizing external architecture risk portraying Gambians and Cubans as passive objects of forces beyond their control, denying the agency any serious historical account requires? Recognizing a chokepoint does not require denying that nations navigate, resist, and at times reproduce their own constraints through domestic choice. It requires insisting that the menu of options available to that navigation is not neutral and was not generated by nature.

A third objection is methodological: this essay has not attempted a quantitative decomposition of how much of either country's strain is attributable to external architecture versus domestic governance, and given how thoroughly the two interact, such a decomposition may not be achievable with the confidence a table of percentages would imply. The claims here should be read as a framework for reallocating attention rather than a settled causal apportionment. What can be said with confidence is directional: the chokepoint contribution is very likely underweighted relative to its actual causal role, not because it is dominant, but because it is currently treated as close to negligible in public discourse when the empirical record, particularly Cuba's documented embargo effects and Gambia's documented debt-service burden, does not support that treatment.

A fourth objection concerns misuse: emphasizing external architecture has real historical precedent as a rhetorical shield used by governments seeking to deflect accountability for their own failures, a concern both the Gambian and Cuban governments' own rhetoric illustrates. The response is not to abandon structural explanation to avoid this misuse, since the same risk attaches to any structural account of any social outcome, but to insist on the specificity that makes misuse harder: naming the exact debt instrument, the exact correspondent-banking decision, the exact embargo provision, rather than gesturing at colonialism or imperialism as an undifferentiated background condition. Specificity is what separates an explanation from an excuse, and it is also the specificity that section eight argues is required to reopen intersubjective recognition in the first place.

8 Reopening Reachability

If material separation, institutional abstraction, collapse, apathy, and political reproduction form a genuine loop, then interventions aimed only at the psychological register, more vivid appeals, more urgent imagery, will remain limited, because they leave the architecture that

regenerates apathy fully intact. Interventions aimed only at the structural register, debt relief, sanctions reform, correspondent-banking transparency, are necessary but insufficient on their own if the public cannot perceive why they matter, since structural reform requires sustained political will that collapse itself makes difficult to sustain.

What the loop actually calls for is intervention at the interface between the psychological and the structural: treating the causal chain from a specific sanctions clause, or a specific correspondent bank's de-risking decision, to a specific blackout as a story told with the same causal specificity that Slovic's research shows is required to generate response to a single identified victim, rather than the abstracted, causeless suffering that dominates most reporting on the Global South. It means presenting luxury consumption and infrastructural deprivation as sharing a supply chain, a shipping lane, a currency system, rather than as unrelated categories of news, because that shared architecture, not any coincidence of timing, is the actual fact obscured by ordinary reporting practice. None of this guarantees that restored recognition produces the political will to dismantle chokepoints whose beneficiaries have every incentive to keep them opaque. The alternative, treating apathy as a fixed baseline to be worked around rather than a produced and reversible condition, forecloses the attempt before it is made.

9 The Discomfort Question Revisited

This essay cannot honestly close without returning to the question that motivated it: how does a person justify spending ten thousand dollars on a vacation, an amount that could substantially alter another person's material circumstances, while people in Gambia and Cuba lack reliable power, food, and medicine? There are serious answers available, and none of them dissolves the underlying asymmetry entirely, which is exactly why the question deserves to remain open rather than resolved too neatly.

One answer appeals to special obligations: people owe more, in the ordinary structure of human moral life, to those with whom they share particular relationships, family, community, national citizenship, than to strangers, and a general obligation to strangers does not straightforwardly override every particular obligation and every legitimate personal project. A second answer appeals to the limits of demandingness: a moral theory that requires every discretionary expenditure to be justified against the maximum possible good it could have done elsewhere is a theory few people can live under without either self-deception or breakdown, and a framework that leaves no room for a person's own life, relationships, and pleasures may defeat itself by being unlivable. A third answer appeals to uncertainty: it is not obvious that redirecting ten thousand dollars toward Gambia's debt service or Cuba's fuel imports would actually reach the people this essay describes rather than being absorbed

by the same chokepoint architecture that produced the deprivation, an uncertainty the philanthropy critique in section four makes considerably more pointed rather than less.

None of these answers, however, addresses the specific form of obligation this essay has been building toward, which is not the general utilitarian question of whether one's resources could do more good elsewhere, the question Peter Singer's drowning-child argument poses so effectively, but a narrower and in some ways more uncomfortable question: what changes when a person recognizes that another's deprivation is not merely known to them in the abstract but produced, in part, by institutions in which they are already a participant, whose dollar-denominated purchases, whose bank accounts, whose consumption of shipping capacity and financial clearing all draw on the same chokepoint architecture that constrains Gambia's debt service and Cuba's fuel imports. This is a question of standing within a shared and already-connected system, not a question of optional charity toward a stranger with whom one has no prior relationship. It does not resolve into a specific dollar figure or a specific required action, and this essay will not pretend that it does. But it does shift the frame from asking how generous one should be to asking what one's relationship, already, is to the architecture producing this outcome, and that shift, modest as it sounds, is where the essay's argument actually terminates: not in a verdict on any individual vacation, but in the claim that the vacation and the blackout were never two separate stories, and that the comfort of believing otherwise is itself a product of the same collapse this essay has been trying to describe.

10 Conclusion

The tourist on the terrace and the family cooking by firewood in Havana are connected by an architecture spanning currency clearance systems, shipping alliances, credit ratings, and sanctions regimes, an architecture whose ordinary operation depends on remaining invisible to the tourist, not through conspiracy but through the routine abstraction every institution operating at that scale requires to function. Effective demand, not human need, is what this architecture was built to register, and everything that follows, meritocratic self-narrative that mistakes survivorship for merit, philanthropy that mistakes the absence of intervention for the relevant baseline, infrastructure evaluation that mistakes one favorable metric for the whole picture, and finally intersubjectivity collapse itself, is a variation on the same underlying failure: mistaking the observed outcome for the full comparison it should have been judged against. Gambia and Cuba arrive at broadly similar deprivation by different roads, one of diffuse debt dependency traceable to a colonial-era monoculture and its structural-adjustment-era successors, one of explicit and enduring embargo entangled with domestic constraint, and both outcomes persist in a world that manifestly has enough power,

food, and medicine to spare. The deepest object connecting effective demand, meritocracy, philanthropy, infrastructure evaluation, and chokepoint capitalism is not wealth or charity or sanctions or debt considered individually, but counterfactual sovereignty: unequal authority over the space of alternatives by which actions, institutions, and distributions come to be declared necessary, successful, generous, efficient, or just. Understanding persistent deprivation as the output of an architecture whose designers, at every level, had the power to choose which counterfactual their success would be judged against, rather than as an unfortunate and unrelated fact about certain unlucky nations, is the necessary first step toward treating it as something other than permanent, and toward asking, honestly, what a person already embedded in that architecture actually owes to the people it has priced out of view.

A Effective Demand versus Need

Let a population of agents $i \in \{1, \dots, n\}$ each have a wealth level $w_i \geq 0$ and a need level $N_i \geq 0$ for some good, where N_i is defined independently of price, for instance as a physiologically or infrastructurally grounded quantity such as minimum caloric intake or minimum reliable electricity supply. Each agent's effective demand at price p is the ordinary budget-constrained optimum,

$$q_i(p) \in \arg \max_{q \geq 0} \{u_i(q) : p q \leq w_i\},$$

not a quantity that collapses to zero once w_i falls short of some unconstrained target; an agent priced out of their preferred bundle still demands whatever smaller quantity their budget allows. A market clears by matching the available supply S against the price at which aggregate demand is feasible,

$$\sum_i q_i(p) \leq S,$$

with no aggregate objective, of the kind a planner might maximize, entering the clearing condition at all: each $q_i(p)$ is determined entirely by agent i 's own (u_i, w_i, p) . Need N_i does not appear anywhere in this condition. Consequently there is no general implication from need to demand,

$$N_i > N_j \not\Rightarrow q_i(p) > q_j(p),$$

and under sufficiently unequal wealth the inequality can run the opposite way from need entirely,

$$N_i > N_j \quad \text{while} \quad q_i(p) < q_j(p),$$

which is the formal statement of the claim that a market can coherently clear while allocating less of a good to the agent who needs it more, whenever wealth and need are weakly or negatively correlated across the population, as between a wealthy tourist's discretionary travel budget and a subsistence household's electricity need. Nothing here requires any agent's demand function to be irrational or any market to be malfunctioning; the divergence between N_i and $q_i(p)$ is exactly what an ordinary competitive market, functioning as designed, produces whenever w_i and N_i diverge.

B The Counterfactual Discipline

A naive evaluation of an action a compares its realized utility $U(a)$ against a single control condition, typically the utility of taking no action at all. This understates the actual structure

of the problem, in which a is drawn from, and should be evaluated against, a feasible set

$$\mathcal{C}_F(a) = \{ c : c \text{ was a reasonably available alternative at decision time } t_0 \},$$

containing every alternative action that could realistically have been taken with the same resources at the time the decision was made, not merely the null action and not merely whatever alternative later turns out, with hindsight, to have performed best. The comparator actually used in evaluation, c_J , should be selected by a rule declared in advance,

$$c_J = \Gamma(\mathcal{C}_F(a), I_{t_0}, K),$$

where I_{t_0} is the information available before the decision and K is a criterion for what counts as a reasonable comparator, fixed independently of how a turns out. The properly attributable value is then

$$V_K(a) = U_K(a) - U_K(c_J).$$

Two failure modes are worth distinguishing. An actor free to select the weakest member of $\mathcal{C}_F(a)$ after the fact, $c^- = \arg \min_{c \in \mathcal{C}_F(a)} U(c)$, can represent almost any a as an improvement, since $V_K(a)$ is then inflated by construction rather than by the merit of a itself. Less obviously, an actor who instead retrospectively selects the single best-performing member of $\mathcal{C}_F(a)$ as the mandatory standard, $c^+ = \arg \max_{c \in \mathcal{C}_F(a)} U(c)$, commits a symmetric error: among several plausible alternatives, one may succeed spectacularly through unpredictable circumstance, and judging every other alternative, including a , a failure relative to that exceptional outcome reproduces the same hindsight problem this appendix is meant to guard against. The discipline this essay is arguing for is neither c^- nor c^+ but c_J : a comparator whose selection rule Γ is justified independently of the observed result, using only information available before the outcome was known. Construction of $\mathcal{C}_F(a)$ and of Γ itself, what alternatives are treated as feasible and by what rule one becomes the standard, is therefore not a neutral analytical preliminary but a further site of power: an actor who influences $\mathcal{C}_F(a)$ and Γ , in addition to choosing a , holds what Section 4.1 calls counterfactual sovereignty, evaluative authority layered on top of ordinary allocative power. What actually guards against c^- - and c^+ -style manipulation in practice is ex ante specification: registering the outcome measures, the comparator, and the decision criterion before $U(a)$ is observed, together with correction for multiple comparisons when several candidate comparators or outcome measures are examined, since significance testing alone, applied after outcomes and comparators have already been chosen freely, offers little protection on its own. This is the formal description of what Sections 4.2 through 4.5 describe informally.

C Survivorship Bias as an Order Statistic

Let n agents each realize an outcome $Y_i = \mu + \varepsilon_i$, where ε_i are independent and identically distributed with mean zero and variance σ^2 , and where no agent-level trait causally affects μ . This describes outcome selection: the winner is

$$i^* = \arg \max_i Y_i,$$

and the maximum of the sample, $Y_{i^*} = \max_i Y_i$, has expectation growing with n even though every individual outcome was generated by an identical, trait-independent process; for ε_i drawn from a light-tailed distribution such as the normal, a standard extreme-value approximation gives

$$E[Y_{i^*}] \approx \mu + \sigma \sqrt{2 \ln n}.$$

For n in the millions, $E[Y_{i^*}]$ substantially exceeds μ , so an extreme winner is expected even though no agent was more capable than any other in the process generating μ .

A second and distinct step, explanation selection, follows once i^* has been identified. Suppose agent i^* is examined for traits $T_1, \dots, T_m$ drawn from some large space of candidate explanations, morning routines, risk postures, personality descriptions, and a trait T_k is reported as the explanation for i^* 's success whenever it appears distinctive relative to the population. This is a multiple-comparisons problem: as m grows, the probability that at least one T_k appears distinctive by chance alone approaches one, independently of whether any T_k actually affects μ , in the same way that testing enough hypotheses against a fixed significance threshold will eventually produce a false positive somewhere in the family even when nothing in the underlying process is true. Outcome selection alone would already make i^* an unrepresentative data point; explanation selection compounds this by searching, after i^* is fixed, for whichever candidate trait happens to correlate with the very outcome that was used to identify i^* in the first place. Establishing that a trait T actually causes elevated μ , rather than merely correlating with the selected winner, requires comparing $P(T \mid \text{success})$ against $P(T \mid \text{no success})$ across the full population of agents, not merely reporting $P(T \mid \text{success})$ for the single winner, which is the quantity meritocratic narrative typically reports, having already passed through both stages of selection without correcting for either.

D The Talent Reserve as Unmeasured Capability

Let $T_i \geq 0$ denote agent i 's latent capability at some task, and let $O_i \geq 0$ denote the institutional opportunity available to i to demonstrate that capability, credentialing access,

employment position, market reach, geographic proximity to whoever is measuring. An institution's recognition function R maps latent capability and opportunity to visible, measured capability,

$$T_i^{\text{visible}} = R(T_i, O_i), \quad R(T_i, 0) = 0,$$

where the boundary condition states only that capability cannot be recognized in the complete absence of any opportunity to exercise it. Crucially, R is increasing in O_i for fixed T_i , so that

$$O_i \approx 0 \implies T_i^{\text{visible}} \approx 0,$$

but this implication does not run in reverse:

$$T_i^{\text{visible}} \approx 0 \not\Rightarrow T_i \approx 0.$$

An institution that reads T_i^{visible} as though it were T_i mistakes a shortage of opportunity for a shortage of capability whenever this second, invalid implication is assumed rather than tested. The aggregate talent reserve is the total capability present in a population but not registered by current recognition procedures,

$$\mathcal{R}_T = \sum_i \left(T_i - T_i^{\text{visible}} \right)_+,$$

using the positive-part operator since T_i^{visible} cannot exceed T_i under a recognition function that does not fabricate capability. Brain drain, in this notation, concerns the departure of agents with high T_i^{visible} to locations offering higher reward for the same recognized capability; the talent reserve concerns agents with $T_i - T_i^{\text{visible}}$ large, capability that remains geographically present but institutionally unregistered, whether because O_i is low or because, as Section 4.3 argues, an agent may rationally suppress the portion of $R(T_i, O_i)$ that would otherwise be visible given available O_i .

E Chokepoint Extraction as a Bottleneck Function

Model global trade as flow through a network in which most nodes have high capacity but a small subset C of edges, the chokepoints, have capacity far below the flow that would otherwise clear. A capacity constraint alone, $f' = \min(f, \kappa_c)$, describes any bottleneck, including a congested bridge that no single actor controls, and does not yet capture what makes a chokepoint politically and economically distinctive. What distinguishes a chokepoint is that passage through it is governed by a controlling actor j , whose decisions x_j , a bank's risk tolerance, a sanctions carve-out, a shipping alliance's routing policy, set both the effective

capacity and the price of access,

$$\kappa_c = \kappa_c(x_j), \quad \tau_c = \tau_c(x_j), \quad \rho_c = \rho_c(x_j) \in \{0, 1\},$$

where ρ_c is a binary compliance or exclusion decision determining whether a given counterparty may use the passage at all. Realized flow for a constrained counterparty is then

$$f' = \rho_c \cdot \min(f, \kappa_c(x_j)),$$

and the controller can extract a rent up to

$$r_{\max} = f' \cdot (\tau_c(x_j) - \text{marginal cost of granting access}),$$

bounded above by the constrained counterparty's willingness to pay rather than by any cost of provision. What separates this from an ordinary bottleneck is the combination of two conditions: the controller's decision has a large effect on the counterparty's realized flow,

$$\left| \frac{\partial f'}{\partial x_j} \right| \gg 0,$$

while the counterparty has little or no ability to route around the constraint through an alternative passage,

$$\text{Substitutability}(c) \approx 0.$$

A congested bridge satisfies the first condition only weakly, since no single actor's discretionary decision governs it, and often fails the second, since alternative routes typically exist even if slower. A reserve-currency clearing system, a dominant shipping alliance, or a sanctions-enforcing correspondent bank typically satisfies both, which is the formal asymmetry underlying Section 5's claim that chokepoint control yields extraction, and leverage, disproportionate to any ordinary measure of market share.

F The Collapse Loop as a Discrete Dynamical System

Let $s_t \in [0, 1]$ denote the degree of material separation, a_t the degree of institutional abstraction, c_t the degree of intersubjective collapse, and p_t the political pressure for reform, all normalized to $[0, 1]$ at time step t . Sections 3 and 8 propose the coupled system

$$a_t = \alpha s_t, \quad c_t = \beta a_t, \quad p_t = \gamma (1 - c_t), \quad s_{t+1} = s_t + \eta (1 - s_t) - \delta p_t s_t,$$

where $\alpha, \beta, \gamma, \delta \in (0, 1)$ are transmission coefficients and $\eta \in [0, 1)$ is an exogenous rate at which chokepoint architecture regenerates separation absent reform pressure. The regeneration term $\eta(1 - s_t)$ pushes s_t toward 1 as separation is regenerated, while the reform term $\delta p_t s_t$ pulls it back toward 0 in proportion to current separation and current political pressure; both terms vanish at the relevant boundary, so $s_{t+1} \in [0, 1]$ whenever $s_t \in [0, 1]$, and the pathological behavior of an earlier, unbounded version of this model does not arise here.

A fixed point s^* satisfies $h(s^*) = 0$ for

$$h(s) = \eta(1 - s) - \delta \gamma (1 - \alpha \beta s) s.$$

Since $h(0) = \eta \geq 0$ and, whenever $\alpha \beta < 1$, $h(1) = -\delta \gamma (1 - \alpha \beta) < 0$, continuity of h guarantees by the intermediate value theorem at least one fixed point $s^* \in (0, 1]$ whenever $\eta > 0$, and $s^* = 0$ is the unique fixed point when $\eta = 0$. Because $\partial h / \partial \eta = 1 - s > 0$ for all $s \in [0, 1)$, increasing the exogenous regeneration rate η shifts h upward at every interior point and, given that h is decreasing in s over the relevant range for modest $\alpha \beta$, moves the equilibrium s^* higher; conversely, increasing δ or γ , the strength of reform response to collapse-reduced pressure, shifts h downward and lowers s^* . This is the formal sense in which Section 8 argues that intervention limited to reducing c_t alone, without also reducing η or strengthening δ , leaves an equilibrium level of separation in place rather than driving s_t to zero.

References

- [1] Banerjee, A. and Duflo, E. *Poor Economics: A Radical Rethinking of the Way to Fight Global Poverty*. New York: PublicAffairs, 2011.
- [2] Boltanski, L. *Distant Suffering: Morality, Media and Politics*. Cambridge: Cambridge University Press, 1999.
- [3] Giblin, R. and Doctorow, C. *Chokepoint Capitalism: How Big Tech and Big Content Captured Creative Labor Markets and How We'll Win Them Back*. Boston: Beacon Press, 2022.
- [4] Levinas, E. *Totality and Infinity: An Essay on Exteriority*. Translated by Alphonso Lingis. Pittsburgh: Duquesne University Press, 1969.
- [5] Sandel, M. *The Tyranny of Merit: What's Become of the Common Good?* New York: Farrar, Straus and Giroux, 2020.
- [6] Singer, P. "Famine, Affluence, and Morality." *Philosophy and Public Affairs* 1, no. 3 (1972): 229–243.
- [7] Slovic, P. "If I Look at the Mass I Will Never Act: Psychic Numbing and Genocide." *Judgment and Decision Making* 2, no. 2 (2007): 79–95.
- [8] Castiglione, B. *The Book of the Courtier*. Translated by George Bull. London: Penguin Classics, 1967 [1528].
- [9] Coface. "Gambia, Republic of the: Economic Studies and Country Risk." Coface Country Risk Files, 2025.
- [10] International Monetary Fund. *The Gambia: 2025 Article IV Consultation*. IMF Country Report No. 25/151. Washington, DC: International Monetary Fund, 2025.
- [11] Government of The Gambia, Ministry of Finance and Economic Affairs. *2024 Annual Macro-Fiscal Report*. Banjul: Ministry of Finance and Economic Affairs, 2025.
- [12] World Bank. "Access to Electricity (% of Population): The Gambia." World Development Indicators, 2023.
- [13] Sangal, A. and Regan, H. "Cuba Plunged into Darkness as Electrical Grid Suffers Fresh Collapse." CNN, October 18, 2024.
- [14] Spectrum staff. "Cuba's Energy Crisis: A Systemic Breakdown." *IEEE Spectrum*, July 2025.

- [15] World Food Programme. "Cuba." WFP Country Page, accessed July 2026.
- [16] United States Department of Agriculture, Economic Research Service. *International Food Security Assessment*. Report No. ERR-340. Washington, DC: USDA ERS, 2024.
- [17] MarketScreener/Reuters. "Cuba on Edge as US Seizure of Oil Tanker Puts Supply at Risk." January 2026.
- [18] NBC News. "Cubans Brace for Even More Economic Devastation amid Threat of No Venezuelan Oil." January 2026.