

The Useless Work Problem in Blockchain Economies: Thermodynamics, Semantic Structure, and Cognitive Control

Flyxion

November 27, 2025

Abstract

Blockchain and cryptocurrency systems have been widely promoted as revolutionary infrastructures for finance, computation, and distributed coordination. Yet the overwhelming majority of deployed blockchain systems fail to produce meaningful civilizational work. They create no new scientific insight, no medical capacity, no epistemic coherence, and no infrastructural resilience. Instead, they propagate speculative assets, synthetic volatility, and illusions of decentralization, attracting global talent into architectures that generate negligible entropy-reducing structure.

This paper develops a comprehensive critique of blockchain economies grounded in a thermodynamic and information-theoretic conception of value, together with tools from self-organizing neural dynamics, predictive coding, sheaf theory, generative modeling, distributed control, and adaptive cognitive optimization. Building on foundational work on irreversibility and useful work in computation (Landauer 1961; Bennett 1973), on the physics of self-organization and replication (Schrödinger 1944; England 2013), on adaptive neural architectures and resonance-based learning (Grossberg 1976; Grossberg 1987; Grossberg 1990), on predictive coding and free-energy–minimizing inference (Friston 2005; Millidge et al. 2021), on cellular sheaves and cohomology (Hansen and Ghrist 2018; Seely 2025), and on flow-map distillation and diffusion-type generative models (Sohl-Dickstein et al. 2015; Song and Ermon 2019; Tong et al. 2025), we argue that blockchains are not merely inefficient but structurally incapable of producing value as entropy-reducing, coherence-expanding work.

We further integrate insights from distributed algorithms and consensus theory (Lynch 1996; Olfati-Saber et al. 2007; Bamieh and Gayme 2012), semantic and categorical infrastructure (Spivak 2014; Ghrist 2014), and cognitive loop optimization (Cheng et al. 2025). These perspectives jointly reveal why blockchain cannot serve as the backbone of any meaningful semantic, infrastructural, or epistemic system, and why redirecting talent and energy away from speculative ledger maintenance toward genuine problem-solving is a matter of civilizational urgency.

Contents

1	Introduction	3
2	Related Work	4
3	Value as Entropy-Reducing Work	5
4	Blockchain as Negative-Value Computation	6
5	Illusion of Decentralization and Simulated Agency	7
6	Cohomological Obstructions to Semantic Coherence	8
7	Generative Mismatch and Flow-Map Distillation	9
8	Semantic Infrastructure and Categorical Coherence	10
9	Integration with Distributed Control and CLIO	10
10	Opportunity Cost and Civilizational Risk	11
11	Conclusion	12
12	Broader Implications	13
13	Policy Recommendations	14
14	Future Work	16
15	Limitations	18
	Appendix A: Thermodynamic Foundations of the Value Functional	20
	Appendix B: Cohomological Obstructions	21
	Appendix C: Teacher–Data Mismatch	23
	Appendix D: Simulated Agency	24
	Appendix E: Categorical No-Go Theorem	25
	Appendix F: Distributed-Control Inefficiency	27
	Appendix G: Dissipative Dynamics	28

Appendix H: Sheaf No-Go for Identity	29
Appendix I: Variational Derivation	30
Appendix J: Token Economies and Market Clearing	32
Appendix K: Epistemic Failure	33
Appendix L: Semantic Verification Complexity	34
Appendix M: Bubble Dynamics	35
Appendix N: Semantic Information Geometry	37

1 Introduction

Blockchain has emerged as one of the most widely publicized computational paradigms of the early twenty-first century. Popular treatments present distributed ledgers as a breakthrough in trustless coordination, censorship resistance, decentralization, and financial autonomy (Narayanan et al. 2016). At the same time, empirical measurements of mining pools, validator dominance, and network participation have repeatedly shown strong structural centralization in major cryptocurrencies such as Bitcoin and Ethereum (Gencer et al. 2018), while transaction-fee dynamics and speculative cycles have been analyzed as a form of market microstructure rather than infrastructural innovation (Easley et al. 2019).

From a broader scientific perspective, blockchain’s promises must be evaluated against what we know about computation, control, and self-organization. Thermodynamic analyses of computation emphasize that useful work must be distinguished from heat dissipation (Landauer 1961; Bennett 1973). Statistical physics of self-organization highlights how certain driven systems exploit entropy gradients to maintain low-entropy configurations and support persistent function (Schrödinger 1944; England 2013). In parallel, decades of work in neural computation and cognitive architectures—from early self-organizing feature detectors and adaptive pattern classification (Grossberg 1976) through competitive learning and adaptive resonance theory (Grossberg 1987; Grossberg 1990) to predictive coding and active inference (Friston 2005; Millidge et al. 2021)—has shown how agents can harness feedback to reduce prediction error and stabilize meaningful internal representations.

Against this background, the central claim of this paper is simple: *most blockchain and cryptocurrency systems do not perform useful work in any thermodynamically or cognitively meaningful sense*. They do not generate new structure in physical, biological, or epistemic domains. They do not improve our capacity to predict or control real systems. Instead, they concentrate energy, hardware, and human talent into maintaining redundant ledgers and

supporting speculative markets.

The argument proceeds in layers. First, we formalize a notion of value as entropy-reducing, coherence-enhancing work, synthesizing insights from thermodynamics, control theory, and self-organizing neural dynamics. Second, we analyze blockchain mechanisms as instances of negative-value computation: systems that expend energy and attention without producing structure. Third, we use recent mathematical developments in sheaf cohomology for predictive coding (Hansen and Ghrist 2018; Seely 2025) and flow-map distillation without data (Tong et al. 2025) to provide explicit models of blockchain failure modes. Finally, we contrast blockchain architectures with alternative directions grounded in distributed control (Lynch 1996; Olfati-Saber et al. 2007), semantic infrastructure (Spivak 2014; Ghrist 2014), and cognitive loop optimization (Grossberg 1987; Grossberg 1990; Cheng et al. 2025).

2 Related Work

A large body of literature examines blockchain systems from perspectives including cryptography, economics, distributed systems, and financial market structure. Much of this work focuses on internal technical properties of blockchains—consensus mechanisms, crypto-economic incentives, security assumptions, and scalability challenges—rather than on their alignment with thermodynamic, semantic, or cognitive criteria for useful work.

From the cryptographic and systems perspective, introductory surveys and monographs such as Narayanan et al. (Narayanan et al. 2016) present blockchains as trustless ledgers designed to resist adversarial modification. These analyses help clarify the cryptographic primitives and data structures underlying major cryptocurrencies. However, they generally take for granted that ledger maintenance constitutes meaningful infrastructural work, without evaluating its thermodynamic or semantic contribution.

Empirical studies of decentralization offer a more critical view. Measurements of mining-pool concentration and validator dominance in Bitcoin and Ethereum have shown persistent structural centralization (Gencer et al. 2018). Analyses of transaction fees and market behavior further reveal that blockchain ecosystems largely follow the dynamics of speculative asset markets rather than serving as efficient computational platforms (Easley et al. 2019). These findings align with the present paper’s claim that most blockchain activity reflects synthetic volatility rather than structural value creation.

Work on distributed systems and consensus algorithms provides an important contrast. Classical analyses of fault-tolerance, synchrony, and multi-agent agreement (Lynch 1996; Olfati-Saber et al. 2007) establish formal criteria for when distributed architectures support robust coordination. Recent studies of power networks and large-scale synchronization (Bamieh and Gayme 2012) highlight that efficient consensus requires careful management of energy flow and communication structure—issues that blockchain protocols rarely address. In this sense, blockchain consensus resembles a degenerate or inefficient instance of distributed

control, magnifying coordination costs without delivering corresponding benefits.

Economic critiques of cryptocurrency markets further underscore this point. The literature on mining incentives, fee markets, and the microstructure of cryptocurrency exchanges shows that token-based economies often behave like highly leveraged speculative arenas, shaped by wealth concentration and market manipulation rather than genuine productivity or informational efficiency. These analyses reinforce the view that blockchain systems primarily channel capital toward arbitrage and rent extraction rather than toward entropy-reducing civilizational work.

In parallel, there is growing interest in applying formal methods—ranging from category theory and sheaf theory (Spivak 2014; Ghrist 2014) to predictive coding and cohomological models of information integration (Hansen and Ghrist 2018; Seely 2025)—to understand how complex systems maintain semantic and structural coherence. These lines of research stand in sharp contrast to blockchain’s purely syntactic consensus mechanisms and help clarify why ledger-based systems struggle to serve as semantically meaningful infrastructures.

Finally, recent work on diffusion models, flow-map distillation, and teacher–student generative dynamics (Sohl-Dickstein et al. 2015; Song and Ermon 2019; Tong et al. 2025) provides tools for understanding how models diverge when trained on pathological or self-referential data distributions. Blockchain systems exemplify such divergence: the ledger evolves based on its own internally generated data, with no mechanism for error-corrective feedback from the external world.

In summary, while existing literature has developed robust analyses of blockchain mechanics, security, incentives, and market behavior, relatively few works consider whether blockchains satisfy the fundamental criteria for value—as entropy-reducing, coherence-enhancing work within real physical, semantic, and cognitive domains. The present paper seeks to fill this gap by integrating thermodynamic, neural, categorical, and generative perspectives to evaluate blockchain’s structural contribution to civilization.

3 Value as Entropy-Reducing Work

A natural starting point is the classical insight that life and organized systems persist by extracting useful work from environmental entropy gradients (Schrödinger 1944). Recent developments in nonequilibrium statistical physics clarify how driven systems can maintain and even increase their structural complexity under suitable conditions (England 2013). In parallel, work on the thermodynamics of computation emphasizes that logic operations have irreducible energetic costs when they erase information (Landauer 1961; Bennett 1973).

At the same time, neural models of self-organization and learning have long emphasized that adaptive systems reduce a form of “internal entropy” by stabilizing useful categories and predictive codes. Grossberg’s early work on adaptive pattern classification and universal recoding (Grossberg 1976) and his later development of competitive learning and adaptive

resonance theory (ART) (Grossberg 1987; Grossberg 1990) provide detailed dynamical accounts of how neural systems can reconcile plasticity and stability: they remain sensitive to new information while preserving established, behaviorally useful representations. In more recent formulations, predictive coding and active inference cast similar ideas in terms of free-energy minimization and hierarchical Bayesian updating (Friston 2005; Millidge et al. 2021).

Let S denote a coarse-grained entropy field over some domain Ω , representing either physical, informational, or semantic uncertainty. A technological system is said to produce *value* when it transforms available free energy into reductions of S that improve the system’s ability to predict, control, or sustain relevant processes. While one can formalize this in various ways, a generic functional

$$\Delta\mathcal{V} = - \int_{\Omega} \partial_t S(x, t) w(x, t) d^3x \quad (1)$$

captures the idea that value corresponds to reducing entropy in regions weighted by w , which reflects their importance for future action.

From the control-theoretic perspective, there are information-theoretic limits on how much a controller can reduce uncertainty given finite channel capacity and energy (Touchette and Lloyd 2000). From the cognitive and neural perspective, adaptive systems must balance learning and stability to avoid catastrophic forgetting while still incorporating new information (Grossberg 1987; Grossberg 1990). In either view, a system that merely consumes energy without producing enduring reductions in S —for example, by performing logically trivial operations or maintaining redundant states—does not generate value in this sense.

This thermodynamic and informational notion of value will be our benchmark for assessing blockchain systems.

4 Blockchain as Negative-Value Computation

The two dominant blockchain consensus mechanisms, Proof of Work (PoW) and Proof of Stake (PoS), instantiate distinct failure modes relative to (1).

In PoW systems, mining nodes repeatedly solve cryptographic puzzles chosen precisely for their uselessness: the computations have no purpose other than demonstrating expenditure. From the standpoint of Landauer’s principle and related analyses of irreversibility in computation (Landauer 1961; Bennett 1973), such systems are paradigmatic examples of heat production without work. The entropy field S of any physically or epistemically relevant domain—medicine, infrastructure, climate, materials—is unchanged. Only the ledger state is updated, and the ledger itself is a passive record.

In PoS systems, mass energy expenditure is replaced by capital locking and staking. While

this reduces direct thermodynamic cost, it still does not contribute to entropy reduction in any meaningful domain. Instead, PoS mechanisms amplify existing wealth distributions and centralize control around large stakeholders, as empirical studies of validator sets and governance participation suggest (Gencer et al. 2018). Again, the ledger changes, but no new structure emerges in the world.

We can express the coherence efficiency of a system as

$$\eta_{\text{coh}} = \frac{E_{\text{struct}}}{E_{\text{in}}},$$

where E_{in} is total energy input and E_{struct} the portion that contributes to enduring structural organization in relevant domains. For blockchain systems, $E_{\text{struct}} \approx 0$, and thus $\eta_{\text{coh}} \approx 0$. In many cases, the opportunity cost of energy and talent diversion makes the effective contribution to $\Delta\mathcal{V}$ strictly negative.

5 Illusion of Decentralization and Simulated Agency

The rhetoric of decentralization suggests that blockchain users participate in genuinely distributed control structures. However, distributed algorithms and networked control theory offer precise notions of consensus, robustness, and cooperation (Lynch 1996; Olfati-Saber et al. 2007). These frameworks are typically concerned with whether local update rules converge to global agreement, how sensitive they are to delays and faults, and what costs are incurred for synchrony (Bamieh and Gayme 2012).

Empirical analyses of major cryptocurrency networks have shown that mining and validation often concentrate in a small number of pools or entities (Gencer et al. 2018). Decision-making—for example, protocol changes or parameter updates—depends on social coordination among a small elite, often using off-chain communication channels. In this sense, decentralization is largely an experiential or symbolic property rather than an operational one.

This gap between perceived and actual control can be understood through the lens of predictive coding and active inference models of cognition (Friston 2005; Millidge et al. 2021). Agents maintain generative models of their environment and act to minimize expected prediction error. Environments that provide dense, salient feedback cues—such as rapidly updating token prices or volatile markets—can create strong subjective impressions of agency even when the underlying causal influence is negligible. Grossberg’s analyses of competitive learning and adaptive resonance (Grossberg 1987; Grossberg 1990) likewise emphasize how feedback signals can stabilize internal codes without necessarily improving environmental control.

Blockchain ecosystems excel at generating such feedback loops. Token holders, traders, and small-scale validators experience continuous updates, graphical interfaces, and pseudo-

participation in governance. Yet the actual levers of protocol-level change and resource allocation remain highly concentrated. The system thus functions as a generator of *simulated agency*: a regime in which agents feel they are acting while their actions have little effect on the underlying dynamics.

In terms of value, such systems redirect cognitive and affective resources toward synthetic control problems, rather than toward domains where action could meaningfully reduce entropy or increase structural coherence.

6 Cohomological Obstructions to Semantic Coherence

To analyze blockchain’s inability to support meaningful semantic or infrastructural roles, it is useful to draw on sheaf theory and cohomology, as applied to information flow and predictive coding. Cellular sheaf models provide a way to represent local data and compatibility conditions across a network, with cohomology classes capturing global obstructions to consistency (Hansen and Ghrist 2018; Ghrist 2014).

Seely’s recent work on predictive coding networks shows how linear predictive coding architectures can be framed as relative cochain complexes, with boundary conditions and residuals encoded in coboundary operators (Seely 2025). In that setting, the relative cohomology group

$$H_{\text{rel}}^1(G; \mathcal{F}) \cong \ker D^\top$$

captures harmonic components of the residual that cannot be eliminated by any amount of local updating. Intuitively, a nontrivial class $[b] \in H_{\text{rel}}^1$ indicates an irreducible contradiction in the global predictions encoded by the network.

From a historical perspective, this sheaf-cohomological formulation can be seen as a mathematically sharp elaboration of challenges already present in earlier neural models. Grossberg’s work on adaptive pattern classification and ART (Grossberg 1976; Grossberg 1987; Grossberg 1990) emphasizes the difficulty of maintaining stable, globally consistent category codes under local learning rules. Seely’s cohomology-based residuals provide a way to formalize such irreducible tensions as global obstructions rather than merely local errors.

Blockchain consensus can be modeled in analogous terms: each node maintains local ledger data, and consensus rules enforce certain compatibility relations. However, because these compatibilities are defined purely syntactically (agreement on block hashes and transaction order) and ignore semantic grounding, there can be persistent global contradictions between the ledger and the real-world states it purports to represent. Forks, reorgs, orphaned blocks, and mismatches between on-chain and off-chain facts all contribute to a kind of *semantic* harmonic residual.

In cohomological language, one can think of the blockchain as supporting a sheaf whose sections are ledger states and whose restriction maps implement consensus rules. The existence

of nontrivial harmonic residuals means that there is no globally consistent assignment of semantic meaning compatible with both the ledger and the external world. No amount of additional computation or block production can resolve this tension, because the obstruction is structural, not quantitative.

This perspective explains why blockchain systems struggle to serve as authoritative registries of real-world assets, identities, or complex contractual relations: their internal consensus mechanisms do not align with the external constraints needed for semantic coherence.

7 Generative Mismatch and Flow-Map Distillation

Recent advances in diffusion-based generative modeling and flow-map distillation provide a complementary lens on blockchain epistemic failure. Score-based and diffusion models learn generative dynamics by estimating gradients of the data distribution or discretizations thereof (Sohl-Dickstein et al. 2015; Song and Ermon 2019). Flow-map methods approximate the time- T map of an underlying continuous-time dynamical system, with distillation techniques seeking to compress teacher trajectories into more efficient student models (Tong et al. 2025).

A central issue identified by Tong et al. is *teacher–data mismatch*: if the student is trained on samples that are not representative of the teacher’s true generative process, the resulting approximations can diverge in systematic ways. In their analysis, the flow map φ_u corresponding to a velocity field u is approximated by a parameterized map f_θ , and mismatches between the distributions used for training and the actual trajectory distribution lead to compounding errors.

Blockchain consensus dynamics can be seen as an extreme case of this phenomenon. The “teacher” process is the actual evolution of the world, with its physical, legal, and social constraints. The “student” is the blockchain ledger, which attempts to encode relevant states and transitions. However, the data used to update the ledger—transactions and blocks—are sampled from a highly constrained, self-referential subset of possible events. They do not reflect the full generative process of the world, nor is there any mechanism to ensure coverage or correction.

In effect, the ledger is a student model trained on a pathological dataset: the interactions of agents already incentivized by the ledger itself. Just as in teacher–data mismatch scenarios (Tong et al. 2025), compounding divergence is inevitable. The ledger’s state drifts farther from any faithful embedding of the real world, yet the consensus mechanism has no access to external gradients that could correct this drift.

From the standpoint of value, such a system cannot reduce entropy in any meaningful sense; it merely evolves according to its own synthetic dynamics, insulated from the generative structure of the domains we care about.

8 Semantic Infrastructure and Categorical Coherence

If blockchains cannot support meaningful semantic roles, what would a more appropriate infrastructure look like? Work in applied category theory and topological data analysis has emphasized the importance of functorial and sheaf-theoretic structures for modeling complex systems of meaning and constraint (Spivak 2014; Ghrist 2014). In these frameworks, objects represent semantic entities or contexts, morphisms represent transformations or inclusions, and functors map between levels of description while preserving structure.

A semantic infrastructure grounded in these principles would place functorial consistency at its core, ensuring that transformations between representations preserve the relational structure that carries semantic significance. It would exhibit a sheaf-like form of locality, allowing compatible local data patches and constraint systems to be glued into coherent global sections while making any resulting obstructions explicitly detectable. Finally, it would maintain a commitment to revisability, permitting its structures to be updated as new information becomes available and thereby reflecting the open-ended and evolving character of scientific, civic, and infrastructural knowledge.

Neural architectures from Grossberg’s ART models (Grossberg 1987; Grossberg 1990) to modern predictive coding networks (Friston 2005; Millidge et al. 2021) provide concrete examples of systems that approximate these desiderata: they maintain hierarchies of representation, exploit feedback loops to reconcile bottom-up and top-down constraints, and continually revise their internal codes in light of new prediction errors. Recent sheaf-theoretic work makes these locality and consistency requirements explicit (Hansen and Ghrist 2018; Seely 2025).

Blockchain systems, with their rigid immutability and purely syntactic consensus rules, are poorly suited to this task. They cannot distinguish between semantically meaningful and meaningless updates, cannot express higher-level constraints that cut across ledger entries, and cannot gracefully accommodate error correction or revision. In categorical language, one might say that their underlying category of states and morphisms is too impoverished to capture the relevant semantics, and their “functor” into the world is neither faithful nor full.

In terms of value, this means that energy and attention invested in blockchains is not building an infrastructure capable of supporting complex, semantically rich systems of coordination. It is building a narrow, brittle substrate optimized for token accounting.

9 Integration with Distributed Control and CLIO

Distributed control theory and multi-agent consensus provide mature frameworks for designing systems that coordinate local agents toward global objectives (Lynch 1996; Olfati-Saber et al. 2007). These frameworks explicitly account for communication constraints, delays, stability, and robustness, and can be applied to domains such as power grids, sensor networks, and

robotic swarms (Bamieh and Gayme 2012). They provide tools for evaluating whether a given architecture actually improves system performance.

By contrast, blockchain consensus provides at best an extremely inefficient broadcast-and-append primitive with no guarantees of optimality, robustness, or semantic alignment. It is rarely the appropriate abstraction for real-world distributed control problems.

The Cognitive Loop via In-Situ Optimization (CLIO) framework (Cheng et al. 2025) offers a complementary perspective on adaptive reasoning and control. CLIO models a cognitive or computational agent as performing local, gradient-based optimization within a loop that directly interacts with the environment, allowing the system to refine its internal models and strategies based on immediate feedback. In this view, useful infrastructures are those that support flexible, context-sensitive optimization and rapid error correction.

CLIO’s in-situ optimization loop extends a long tradition of adaptive control in neural systems, including Grossberg’s models of competitive learning, stable category formation, and resonance-based code selection (Grossberg 1987; Grossberg 1990). In these architectures, feedback loops between bottom-up and top-down signals serve to stabilize representations that are both predictive and behaviorally useful, while rejecting inconsistent or noisy patterns. CLIO generalizes this idea to scientific and computational reasoning, emphasizing that effective control requires architectures that can *learn in the loop* while exposed to real-world dynamics.

Blockchain architectures, with their long finality times, immutability, and heavy reliance on global consensus, are almost the opposite of CLIO-compatible substrates. They cannot rapidly adapt to local conditions, cannot be easily rolled back when errors are detected, and cannot represent the nuanced gradients needed for in-situ optimization. As such, they are ill-suited for supporting the kinds of self-adaptive reasoning and control that scientific and infrastructural systems require.

10 Opportunity Cost and Civilizational Risk

Finally, there is the question of opportunity cost. The computational and human resources devoted to blockchain systems are nontrivial: large-scale mining operations, specialized hardware, and substantial pools of technical talent are engaged in maintaining and extending these networks. From the perspective of value as entropy-reducing work, we must ask what these resources could be doing instead.

The same expertise in distributed systems and cryptography could be applied to secure critical infrastructure, design resilient communication protocols, or build transparent and verifiable systems for scientific data sharing. The same hardware could be used for simulations of climate, materials, or biological systems. The same incentive-design efforts could be directed toward mechanisms that reward genuine contributions to public goods rather than speculative trading on synthetic assets.

In cognitive and neural terms, the analogy is to a brain that devotes most of its plasticity and metabolic resources to maintaining internally generated hallucinations rather than tracking and controlling the external environment. Grossberg’s models of self-organizing neural systems (Grossberg 1976; Grossberg 1990) and predictive coding accounts of cortical function (Friston 2005; Millidge et al. 2021) both emphasize that adaptive systems must align their internal codes with external regularities in order to survive and perform useful work. A blockchain-dominated ecosystem instead aligns internal dynamics with its own speculative feedback loops, drifting farther from the generative processes that matter.

In this light, blockchain is not merely neutral. It is a *semantic and thermodynamic sink*: a system that absorbs energy, attention, and ingenuity while providing little or no contribution to the reduction of entropy in the domains that matter for civilizational stability and flourishing.

11 Conclusion

Drawing on thermodynamics of computation, statistical physics of self-organization, self-organizing neural dynamics, predictive coding, sheaf-theoretic analyses of information flow, flow-map generative modeling, distributed control theory, and cognitive loop optimization, we have argued that most blockchain and cryptocurrency systems do not perform meaningful work in any physically or informationally robust sense. They consume energy without reducing relevant entropy, create illusions of decentralization and agency without corresponding structural control, and lack the semantic and categorical richness required of a genuine infrastructural substrate.

Alternative directions grounded in semantic coherence, revisable structures, and in-situ optimization exist and are already under active development in multiple fields. Neural and cognitive models from Grossberg’s adaptive resonance theory (Grossberg 1987; Grossberg 1990) to modern predictive coding (Friston 2005; Millidge et al. 2021), sheaf-theoretic approaches to information integration (Hansen and Ghrist 2018; Seely 2025), flow-map generative modeling (Tong et al. 2025), and CLIO-style cognitive loops (Cheng et al. 2025) all point toward architectures that use feedback to reduce uncertainty and support interpretable control.

The central challenge is not to find ways to graft these ideas onto blockchain, but to redirect energy and attention away from speculative ledger maintenance and toward infrastructures that actually expand our capacity to understand, predict, and repair the world. A civilization that takes seriously the thermodynamics of computation, the dynamics of learning, and the structure of meaning cannot afford to devote its best resources to systems optimized primarily for synthetic scarcity and volatility.

12 Broader Implications

The analysis in this paper has implications that extend well beyond blockchain systems themselves. At stake is the broader question of how digital infrastructure should interact with physical systems, epistemic institutions, economic structures, and cognitive architectures. The critique developed here points toward several overarching themes relevant to the future of information technology and societal organization.

Semantic Fragility of Digital Institutions

A key implication of the categorical and sheaf-theoretic analysis is that digital infrastructures lacking semantic grounding are structurally incapable of supporting large-scale institutional functions. Systems that encode only syntactic constraints—even if cryptographically secure—cannot maintain the coherence required for identity, governance, or public knowledge. This issue is not unique to blockchain but reflects a broader technological trend toward protocols and platforms that lack mechanisms for semantic integration or error-corrective embedding in real-world constraints.

As digital systems increasingly mediate critical societal functions, the absence of semantic grounding risks creating brittle infrastructures: platforms that can enforce internal formal consistency while drifting arbitrarily far from the real-world structures they purport to represent.

Thermodynamics of Civilizational Computation

The thermodynamic critique raises questions about the allocation of computational resources at the civilizational scale. Most real-world challenges—climate modeling, materials science, biomedical R&D, infrastructure resilience—require massive computation devoted to entropy reduction in high-dimensional state spaces. Systems like Proof-of-Work, which expend energy without producing structure, represent not only inefficiency but negative externality: an opportunity cost measured in lost problem-solving capacity.

This suggests an urgent need to evaluate computational technologies not merely in terms of throughput or decentralization but in terms of their contribution to global entropy reduction and structural coherence.

Epistemic and Cognitive Consequences

The notion of *simulated agency* developed in the main text has broader consequences for digital cognition and collective intelligence. Systems that provide dense feedback without corresponding causal influence can distort user perception, induce addictive or compulsive engagement, and create the illusion of political participation or economic empowerment while

offering no actual control. This phenomenon, which draws on principles from predictive coding (Friston 2005; Millidge et al. 2021) and stability–plasticity dynamics (Grossberg 1987; Grossberg 1990), generalizes to many modern platforms.

Blockchain ecosystems therefore provide a case study in how digital infrastructures can amplify misaligned cognitive loops, divert attention from real-world challenges, and erode epistemic norms.

Implications for Distributed Systems and AI Alignment

The failure of blockchain to produce coherent global states highlights deeper issues in distributed cognition and AI alignment. Systems that cannot reconcile local and global constraints, or that operate without semantic grounding, will inevitably drift. Future large-scale AI systems will require architectures far richer than ledgers: architectures capable of maintaining coherence across levels of abstraction, integrating error-corrective feedback, and updating representations *in situ* (Cheng et al. 2025).

The design principles for aligned AI, resilient infrastructure, and coherent institutions may therefore converge on similar mathematical structures: hierarchical predictive coding, semantic sheaves, and category-theoretic constraint systems.

13 Policy Recommendations

Although this paper focuses on theoretical analysis, several policy implications follow directly from the thermodynamic, semantic, and cognitive critiques of blockchain systems. These recommendations target governments, research institutions, regulatory bodies, and funding agencies.

Reallocate Computational Resources Toward Useful Work

Governments and research institutions should prioritize compute subsidies, research grants, and energy infrastructure for systems that contribute to real entropy reduction—such as climate simulations, biomedical modeling, materials discovery, and infrastructure optimization. Mining-based cryptocurrencies should be excluded from subsidies and may warrant explicit carbon pricing or targeted energy taxation due to their lack of structural output.

Establish Thermodynamic Efficiency Benchmarks

Digital infrastructures should be evaluated according to the extent to which they contribute to structural coherence rather than according to measures such as throughput or nominal decentralization. To support this shift, regulators and standards bodies could adopt quantitative metrics that reflect thermodynamic and semantic performance, including the rate at

which a system reduces entropy in its target domain, indices of semantic consistency derived from sheaf cohomology, and ratios that capture the efficiency with which energy is converted into meaningful structural output. Establishing such benchmarks would allow policymakers and engineers to distinguish systems that genuinely enhance scientific, economic, and civic functioning from those that merely coordinate syntactic data without producing substantive informational or structural value.

Mandate Semantic Consistency for Critical Digital Registries

For systems that manage identity, land records, regulatory compliance, or other forms of critical infrastructure, semantic grounding must be treated as a non-negotiable legal requirement. In practice, this means that digital registries must maintain explicit and verifiable linkages to the real-world constraints they encode, must incorporate auditable mechanisms for detecting and correcting errors, and must employ data structures that remain revisable so that the registry can track evolving physical, legal, and institutional facts. Immutable ledgers such as blockchains are therefore unsuitable for applications in which alignment with physical or legal reality is essential, since their syntactic immutability prevents the continuous semantic reconciliation required by these domains.

Discourage Speculative Token Economies

Because token-based economies are structurally incapable of achieving meaningful market-clearing equilibria, as demonstrated in Appendix J, policymakers should classify cryptocurrency markets as highly leveraged speculative environments rather than as functional financial systems. This classification warrants strong consumer-protection measures, rigorous transparency requirements for exchanges and validator networks, careful limits on retail exposure to high-risk token instruments, and explicit disclosures regarding the systemic volatility and bubble dynamics intrinsic to these markets. Treating token economies in this manner acknowledges their speculative character and mitigates the risks they pose to uninformed participants and to broader economic stability.

Support Research in Semantic and Cognitive Infrastructure

Public investment should be directed toward mathematical and computational frameworks that promote semantic consistency, adaptive control, and the alignment of digital representations with real-world structure. This includes advancing applied category theory, developing sheaf-theoretic methods for information integration, strengthening the theory and practice of distributed control systems, expanding hierarchical predictive-coding architectures, and supporting research on cognitive-loop-based reasoning systems (Cheng et al. 2025). Collectively, these areas form the conceptual foundation for future infrastructures that maintain coherence

across levels of abstraction and remain sensitive to the physical and epistemic constraints of the environments in which they operate.

Promote Civic and Scientific Compute Platforms

The analysis presented in this paper also supports the creation of civic computing platforms modeled on scientific infrastructure rather than on speculative token economies. Such platforms would include national compute commons dedicated to climate modeling, biomedical simulation, and other forms of high-impact scientific inquiry; publicly governed scientific data repositories equipped with explicit semantic guarantees; and sustained funding for the development of open-source semantic middleware capable of supporting large-scale collaborative reasoning. They would further benefit from international coordination aimed at establishing thermodynamic and epistemic standards for AI systems. These measures would redirect computational and intellectual resources toward domains in which entropy reduction carries the highest civic and biological significance, thereby aligning infrastructure investment with long-term societal goals rather than with short-term speculative dynamics.

14 Future Work

The critique presented in this paper does not merely identify limitations of existing blockchain systems; it suggests a broader research direction aimed at constructing infrastructures that produce genuine value in the thermodynamic, semantic, and cognitive sense. Several lines of future work follow from the framework developed here.

Semantic Infrastructures Beyond Ledgers

The categorical and sheaf-theoretic arguments indicate that future cyber-physical infrastructures must be capable of representing semantic constraints, updating in a manner consistent with real-world dynamics, and supporting meaningful error-correction. Developing such infrastructures requires integrating methods from applied category theory (Spivak 2014), topological data analysis (Ghrist 2014), and sheaf-driven predictive coding (Hansen and Christ 2018; Seely 2025).

A promising direction lies in the construction of “semantic middleware” layers that can glue locally generated data clouds into globally coherent structures, detect and resolve cohomological obstructions that arise from incompatible local descriptions, and maintain explicit higher-order constraints needed for stable interpretation across contexts. Systems built along these lines would resemble adaptive neural architectures rather than ledgers, exhibiting the capacity to integrate heterogeneous information sources, correct inconsistencies, and preserve global structure through continual adjustment.

Thermodynamically Efficient Consensus and Control

The distributed-control analyses developed in this paper indicate that consensus and co-ordination algorithms should be generalized to exploit Lyapunov structures, energy-aware routing, and local smoothing if they are to serve as principled foundations for future cyber-infrastructure. A promising research direction involves the construction of hybrid architectures that combine the mathematical discipline of multi-agent consensus dynamics (Olfati-Saber et al. 2007), the stabilizing mechanisms of adaptive control, and the hierarchical feedback characteristics found in predictive- coding models of neural computation (Friston 2005; Millidge et al. 2021). In such architectures, the objective would not be merely to maintain syntactic agreement across agents but to minimize global prediction error or entropy production, thereby aligning distributed computation with the deeper thermodynamic and informational constraints of the environment.

Integration with Cognitive Loop Optimization

The CLIO framework (Cheng et al. 2025) offers a unifying principle for civilizational computation: systems must revise their internal representations while interacting *in situ* with the domains they are designed to serve. A fruitful avenue for future research would involve developing infrastructures that are explicitly compatible with this cognitive-loop perspective, enabling round-trip feedback with real-world data, incorporating mechanisms for semantic revision during operation, and providing access to local gradients that support optimization at scale. Architectures designed along these lines would function less like static ledgers and more like adaptive cognitive systems, capable of maintaining coherence through continual interaction, correction, and refinement.

Thermodynamic Benchmarks for Digital Infrastructure

The value functional derived in this paper indicates the need for a new class of benchmarks for digital infrastructure, not measured by throughput, latency, or transaction finality, but by the extent to which a system contributes to entropy reduction in the domain it intends to support. This motivates the development of metrics analogous to Fisher information, variational free-energy, or Lyapunov functionals, each of which captures a system’s ability to generate and maintain real structural coherence.

A broader research agenda would involve applying such thermodynamic and information-theoretic metrics across a spectrum of domains in order to assess the coherence-producing capabilities of scientific computing platforms, civic-infrastructure systems, AI-assisted reasoning architectures, and collaborative knowledge networks. The goal is to evaluate whether a given digital system provides genuine epistemic or structural value rather than merely coordinating syntactic data flows.

Replacing Token-Based Incentives

The economic analyses presented in the appendices imply that token-based incentive systems, by their very construction, tend to generate bubbles, instability, and volatility rather than productive forms of coordination or work. A forward-looking research program should therefore investigate incentive mechanisms that reward contributions to public knowledge, that reflect thermodynamic measures of entropy reduction, that enhance civic robustness, and that incorporate long-horizon forecasting accuracy as a core principle of value creation.

Such an approach would shift the design of digital infrastructure away from speculation-driven token economies and toward systems that are thermodynamically grounded, semantically coherent, and cognitively adaptive, thereby aligning technological development with the long-term informational and structural needs of society.

15 Limitations

Although this paper integrates thermodynamics, cognitive science, distributed control, category theory, and generative modeling to critique blockchain architectures, several limitations remain that motivate further refinement.

Scope of Critique

First, the analysis focuses on mainstream blockchain architectures—primarily PoW and PoS systems—and widely deployed public ledgers. It does not fully address less common approaches, such as DAG-based ledgers, trusted execution environments, or zero-knowledge rollup systems. While the general conclusions likely extend to these architectures, a fully domain-specific analysis remains to be done.

Thermodynamic Abstraction

The entropy functional developed in the main text is general and intentionally abstract. While it accords with nonequilibrium statistical physics (Schrödinger 1944; England 2013) and with foundational computational thermodynamics (Landauer 1961; Bennett 1973), a more detailed accounting of domain-specific entropy flows (e.g. in climate modeling, biomedical research, or logistics optimization) would sharpen the quantitative arguments.

Partial Treatment of Social and Institutional Constraints

Real-world economic and legal systems impose constraints not captured in the formal models used here. While the categorical and sheaf-theoretic no-go theorems demonstrate structural incompatibility between blockchain syntax and semantic grounding, a more comprehensive

institutional analysis would help clarify the precise boundary conditions under which such systems fail.

Modeling Assumptions in Economic Critiques

The economic critiques rely on standard microeconomic and stochastic models. While these models capture well-known instabilities and non-clearing behavior in speculative markets, they do not fully capture behavioral phenomena, regulatory interventions, or macroeconomic shocks. Extending the analysis to agent-based economic simulations or empirical econometrics would further validate the conclusions.

Generative-Model Analogy Limitations

The teacher–data mismatch analogy (Tong et al. 2025) is precise at the level of distributional divergence but abstracts away from the rich structure of real-world dynamical systems. A stronger formal connection would require specifying the teacher and student processes at a deeper level—possibly using stochastic differential geometry or higher-order variational flows.

Absence of Empirical Case Studies

This paper focuses primarily on theoretical and mathematical critique. Although many empirical studies support the central claims—including measurements of network centralization (Gencer et al. 2018) and market instability (Easley et al. 2019)—a more detailed evaluation of specific blockchain deployments (e.g. supply-chain tracing or on-chain identity systems) would strengthen the applied dimension.

Open Questions in Semantic Infrastructure Design

Finally, while the paper sketches an alternative vision for semantic and thermodynamic infrastructure, it does not provide an explicit construction of an operational system. The tools for such a system—category theory, sheaf theory, predictive coding, and distributed control—are well-developed, but the synthesis remains a conceptual framework rather than a fully implemented architecture.

These limitations are not weaknesses but rather opportunities for future work, reflecting the ambition of grounding digital infrastructure in principled, thermodynamic, semantic, and cognitive theory.

Appendix A: Thermodynamic Foundations of the Value Functional

This appendix derives Equation (1) from first principles in nonequilibrium statistical mechanics and the thermodynamics of computation.

A.1 Entropy, Free Energy, and Useful Work

Let \mathcal{S} denote the Shannon–Gibbs entropy of a coarse-grained field over a domain Ω . In a nonequilibrium steady state, the free energy F of a system satisfies

$$F = \langle E \rangle - TS.$$

A reduction in S at fixed energy represents an increase in free energy available for work:

$$dF = -TdS.$$

Work-producing systems exploit entropy gradients to maintain organization (Schrödinger 1944; England 2013). Following Landauer and Bennett (Landauer 1961; Bennett 1973), any reduction in configurational entropy requires expenditure of physical energy, and the ratio of entropy reduction to input energy quantifies efficiency.

A.2 Deriving the Value Functional

Suppose an engineered system influences an entropy field $S(x, t)$ via dynamics

$$\partial_t S = -\Phi(x, t) + \Xi(x, t)$$

where Φ is the entropy flux generated by the system’s operations and Ξ is exogenous noise.

Let $w(x, t)$ be a weight function representing relevance to actionable goals, similar to attention-weighting in cognitive systems or energy-weighting in nonuniform thermodynamic environments.

Define total useful work done over a time interval $[t_0, t_1]$ by:

$$\mathcal{V} = \int_{t_0}^{t_1} \int_{\Omega} \Phi(x, t) w(x, t) d^3x dt.$$

Using $\Phi = -\partial_t S + \Xi$, and ignoring exogenous noise (no engineered system can claim credit

for entropy reductions arising from exogenous fluctuations),

$$\mathcal{V} = - \int_{t_0}^{t_1} \int_{\Omega} \partial_t S(x, t) w(x, t) d^3x dt.$$

This yields the functional used in the main text.

A.3 Why Blockchain Has $\mathcal{V} \approx 0$

Since PoW and PoS systems modify only ledger entropy—a trivial symbolic subsystem detached from any physical or epistemic gradient—the effective $\partial_t S_{\text{real}}$ is zero. Thus, blockchain energetics satisfy:

$$\mathcal{V}_{\text{blockchain}} \approx 0, \quad \eta_{\text{coh}} = \frac{E_{\text{struct}}}{E_{\text{in}}} \approx 0.$$

Energy expenditure produces heat but no corresponding reduction in meaningful entropy.

Appendix B: Cohomological Obstructions to Blockchain Semantic Coherence

This appendix provides formal derivations for the argument that blockchain systems cannot represent consistent global semantic states because their constraints are purely syntactic and lack the local-to-global compatibility conditions required of a sheaf.

B.1 Sheaves, Restrictions, and Compatibility

Let G be a graph or cell complex indexing local contexts. A sheaf \mathcal{F} assigns to each vertex v a data space $\mathcal{F}(v)$, and to each edge $e = (v, w)$ a restriction map

$$\rho_{v \rightarrow e} : \mathcal{F}(v) \rightarrow \mathcal{F}(e), \quad \rho_{w \rightarrow e} : \mathcal{F}(w) \rightarrow \mathcal{F}(e).$$

A global section $s \in \mathcal{F}(G)$ satisfies

$$\rho_{v \rightarrow e}(s(v)) = \rho_{w \rightarrow e}(s(w)) \quad \forall e = (v, w).$$

In predictive-coding networks, Seely (Seely 2025) models prediction errors as cochains:

$$c \in C^1(G; \mathcal{F}), \quad (Dc)(f) = \sum_{e \in \partial f} \sigma_{fe} c(e),$$

where D is a coboundary operator.

B.2 Harmonic Residuals

Residual error b is harmonic if

$$D^\top b = 0, \quad b \neq Dz \quad \text{for any } z \in C^0.$$

Thus $[b] \in H^1(G; \mathcal{F})$ is a nontrivial cohomology class.

This constitutes an obstruction to a globally consistent state.

B.3 Why Blockchain Cannot Eliminate Harmonic Residuals

Let $\mathcal{F}_{\text{chain}}$ denote the “ledger sheaf,” in which each vertex of the underlying network corresponds to a node holding a local copy of the ledger state, and each edge imposes the requirement that neighboring nodes agree on quantities such as block hash, block height, and transaction order. These conditions ensure only that adjacent ledger views remain syntactically compatible.

The essential difficulty is that the restriction maps in $\mathcal{F}_{\text{chain}}$ govern syntactic coherence alone. They do not encode, enforce, or even represent the semantic relationships that arise in real domains, such as ownership of physical assets, the identity of participants, or the terms and conditions of legal or contractual obligations. As a consequence, blockchain consensus eliminates only those discrepancies that can be expressed as differences in raw ledger data, while leaving untouched the broader class of semantic inconsistencies that manifest as harmonic residuals in a sheaf-theoretic analysis. Such residuals arise precisely because the ledger’s compatibility conditions do not reflect the true structure of the world it attempts to model.

Thus the real-world consistency sheaf $\mathcal{F}_{\text{world}}$ and the ledger sheaf $\mathcal{F}_{\text{chain}}$ are not naturally related by a morphism:

$$\mathcal{F}_{\text{chain}} \not\simeq \mathcal{F}_{\text{world}}, \quad \text{Hom}(\mathcal{F}_{\text{chain}}, \mathcal{F}_{\text{world}}) = \emptyset.$$

Hence there exists a harmonic residual

$$[b] \in H^1(G; \mathcal{F}_{\text{chain}})$$

corresponding to irreconcilable mismatch.

No amount of further block production can reduce this residual: it is a *topological obstruction*, not a quantitative error.

Appendix C: Teacher–Data Mismatch and Consensus Divergence

We formalize the analogy between flow-map distillation (Tong et al. 2025) and blockchain consensus.

C.1 Flow Maps and Distillation

Let $u(x)$ be a velocity field generating flows via

$$\frac{d}{dt}x(t) = u(x(t)).$$

The time- T flow map φ_u satisfies

$$\varphi_u(x_0) = x(T).$$

Distillation approximates φ_u by a parametric student model f_θ trained on samples $(x_0, \varphi_u(x_0))$ drawn from some distribution μ .

If the true state evolves under a distribution $\nu \neq \mu$, then the student accumulates bias:

$$\mathbb{E}_{x \sim \nu}[\|f_\theta(x) - \varphi_u(x)\|^2] = \text{irreducible mismatch error.}$$

C.2 Blockchain as a Self-Referential Student

Let the teacher distribution be given by the real-world state transitions $x_{t+1} = \Phi(x_t)$, which encode the lawful evolution of the underlying physical, legal, or economic system. In contrast, the blockchain operates as a self-updating student whose state evolves according to the transition rule

$$\ell_{t+1} = \Psi(\ell_t, \text{transactions}_t),$$

where the function Ψ represents the consensus mechanism and the internal rules of the ledger.

A crucial structural asymmetry arises from the fact that the distribution of transactions driving these updates is not exogenous. Instead, the transaction distribution μ_{chain} is generated by the ledger itself through the behavioral incentives, constraints, and opportunities encoded in its current state. Formally,

$$\mu_{\text{chain}} = \text{Law}(\text{actions} \mid \ell_t),$$

which expresses the fact that agents sample their actions conditional on the ledger’s own configuration. Thus the blockchain learns only from data it generates, inherits no external

ground truth, and becomes a self-referential student trained on its own outputs rather than on the dynamics of the world it purports to represent.

Thus $\mu_{\text{chain}} \neq \nu_{\text{world}}$ generically.

Consensus enforces agreement within the student model, but not between student and teacher.

Hence divergence:

$$\|\ell_t - \Phi^t(x_0)\| \rightarrow \infty$$

in semantic dimensions.

This formalizes why blockchains become increasingly detached from real-world states and why syntactic consensus cannot repair semantic drift.

Appendix D: Simulated Agency and Stability–Plasticity Dynamics

Here we derive the formal analogy between simulated agency in blockchain ecosystems and well-studied failures in adaptive neural systems, drawing on Grossberg’s work on competitive learning and ART (Grossberg 1976; Grossberg 1987; Grossberg 1990).

D.1 Adaptive Resonance Model

Let I denote input patterns and C stored category representations. Resonance requires:

$$\|I - C_j\| < \rho,$$

where ρ is the vigilance parameter.

If ρ is too low, the system admits spurious categories; if too high, it discards legitimate novelty (catastrophic brittleness).

D.2 Blockchain Feedback Dynamics

Let A_t denote the action-space of users (trading, staking, governance). Let F_t be the feedback signals (prices, block rewards, governance UI).

We define the perceived agency at time t by

$$\mathcal{A}_t = \mathbb{E}[\text{Impact}(A_t) \mid F_t],$$

where A_t denotes the action taken by an agent and F_t represents the information available to the agent at that time. In speculative ecosystems, the information flow F_t is typically dense, rapidly fluctuating, and highly stimulating, creating the appearance of continual

influence and opportunity. However, the actual impact of any individual action on the protocol-level structure is negligible, since the underlying system is driven by global market dynamics, validator incentives, and exogenous speculative forces rather than by local agent interventions. This mismatch between informational richness and structural impotence gives rise to a persistent gap between perceived agency and real causal influence.

Thus:

$$\mathcal{A}_t \gg \text{RealAgency}(A_t).$$

This corresponds to a low-vigilance regime in which internal resonance stabilizes patterns unsupported by external structure:

$$\text{Internal feedback coherence} \not\Rightarrow \text{external control}.$$

D.3 Consequence: Persistent Hallucinated Control Loops

The system’s reinforcement dynamics reward repetitive engagement without corresponding environmental influence—exactly the hallmark of stability–plasticity mismatch in neural models.

Appendix E: A Categorical No-Go Theorem for Semantic Integration

This appendix provides a rigorous categorical argument showing that blockchain systems cannot support a semantics-preserving mapping from real-world state categories into ledger-state categories.

E.1 Categories of Real-World States and Ledger States

Let \mathcal{W} denote the category of real-world states relevant to some domain, such as identity, ownership, or contractual relations. In this category, the objects are world-states w , and the morphisms $\phi : w \rightarrow w'$ represent lawful transitions from one state to another, reflecting the physical, legal, or institutional dynamics that govern the domain.

Similarly, let \mathcal{L} denote the category of blockchain ledger states. Here, the objects are ledger configurations ℓ , and the morphisms $\psi : \ell \rightarrow \ell'$ consist of block extensions, reorganizations, or other protocol-level transitions that update the state of the ledger. These transitions are purely syntactic, reflecting the internal rules of the consensus protocol rather than any external semantic constraints.

E.2 Requirement for Semantic Integration

A semantics-preserving integration of real-world states into ledger states would require the existence of a functor

$$F : \mathcal{W} \longrightarrow \mathcal{L},$$

mapping each world-state to a corresponding ledger-state and each lawful transition in the world to an appropriate ledger transition. For such a functor to serve as a faithful semantic embedding, it must distinguish between distinct world-states, capture the full range of meaningful ledger transitions that correspond to real-world changes, and preserve the commutativity of diagrams encoding legal and physical constraints. In categorical terms, this means that if a composite morphism $\phi \circ \psi$ represents the sequential execution of two lawful world transitions, then the image under F must satisfy

$$F(\phi \circ \psi) = F(\phi) \circ F(\psi),$$

ensuring that the ledger-level representation mirrors the structure of the underlying world.

E.3 No-Go Result

Theorem E.1 (Semantic No-Go) There exists no faithful, full, constraint-preserving functor from \mathcal{W} to \mathcal{L} for any domain in which world-state transitions depend on physical, social, or legal structure not expressible in the ledger's syntax.

Sketch.

Ledger states encode:

$$\ell = (\text{balances, signatures, blockhashes, ...})$$

but not physical/location/legal conditions (identity, possession, contractual satisfaction).

Thus, distinct world morphisms

$$w_1 \xrightarrow{\phi} w_2, \quad w_1 \xrightarrow{\phi'} w_2$$

which differ semantically but not in ledger syntax must map to the same ledger morphism:

$$F(\phi) = F(\phi'),$$

violating faithfulness.

Similarly, many ledger morphisms (reorgs, rollbacks, timestamp manipulations) do not correspond to any lawful world transition, violating fullness.

Finally, real-world diagrams of constraints (physical laws, legal chains of custody, identity verification) do not commute with ledger operations because ledger transitions can occur

without satisfying external constraints.

Thus no such F exists.

This completes the proof.

Appendix F: Distributed-Control Inefficiency of Blockchain Consensus

This appendix shows that blockchain consensus mechanisms are strictly dominated by classical distributed control algorithms in terms of communication overhead, energy cost, and convergence behavior.

F.1 Classical Consensus Model

In multi-agent consensus (Olfati-Saber et al. 2007), agents update:

$$x_i(t+1) = \sum_{j \in N(i)} a_{ij} x_j(t),$$

with weights satisfying stochasticity and connectivity.

Convergence is exponential if the graph is connected.

Communication complexity per round is $O(|E|)$.

F.2 Blockchain Consensus

In PoW/PoS:

$$\ell_{t+1} = \Psi(\ell_t, B_t),$$

with B_t a globally broadcast block.

Communication complexity is:

$$O(n^2)$$

for n nodes (full broadcast and retransmission).

Finality delay is on the order of minutes to hours.

F.3 Inefficiency Theorem

Theorem F.1 (Dominance of Classical Consensus). For any network consisting of n agents and for any task that requires those agents to reach agreement on a shared value, classical distributed-consensus algorithms strictly outperform blockchain-based consensus mechanisms with respect to communication complexity, overall energy expenditure, convergence rate, stability under delay, and robustness to faults.

Sketch. The reason for this dominance is structural. Blockchain protocols require that every update be propagated through global broadcast rather than through localized neighbor interactions, thereby incurring far greater communication costs. They rely on energy-intensive block production, introduce probabilistic rather than deterministic finality, and inevitably permit forks and reorganizations that disrupt convergence. They offer no mechanism for weighted averaging or state smoothing, both of which are central to efficient consensus dynamics, and they lack robustness in the presence of asynchronous delays or variable network conditions.

Taken together, these properties show that blockchain consensus constitutes a strictly dominated strategy within the design space of multi-agent consensus systems: it is more expensive, slower, less stable, and less robust than classical alternatives, while providing no compensatory advantage in terms of control or coordination.

Appendix G: Proof-of-Work as a Dissipative Dynamical System

Here we provide a Lyapunov analysis showing that Proof-of-Work (PoW) mining is a dissipative system that converges to a limit cycle with no useful output.

G.1 Mining as a Dynamical System

Let $h(t)$ denote global hash rate, $d(t)$ difficulty, $r(t)$ rewards.

Mining profitability $\pi(t)$ satisfies:

$$\pi(t) = r(t) - ch(t),$$

where c is cost per hash.

Difficulty adjusts as:

$$\frac{d}{dt}d(t) = \alpha(\text{target} - T(h(t), d(t))),$$

where T is expected block time.

G.2 Lyapunov Functional

Define:

$$\mathcal{L}(h, d) = \frac{1}{2}(T(h, d) - \text{target})^2.$$

Differentiating along trajectories:

$$\dot{\mathcal{L}} = (T - \text{target}) \left(\frac{\partial T}{\partial h} \dot{h} + \frac{\partial T}{\partial d} \dot{d} \right).$$

Given the form of \dot{d} ,

$$\dot{\mathcal{L}} = -(T - \text{target})^2 \leq 0.$$

Thus PoW has a Lyapunov-stable manifold defined by $T = \text{target}$.

G.3 No Useful Work Property

On the manifold, the system settles into a limit cycle where h grows until profit approaches zero.

Entropy produced:

$$\dot{S}_{\text{heat}} = ch.$$

Useful work:

$$\dot{W}_{\text{useful}} = 0.$$

Thus PoW is a stable dissipative system with no useful output.

Appendix H: Sheaf-Theoretic No-Go for On-Chain Identity

This appendix formalizes why blockchain cannot reliably represent real-world identity or physical asset ownership.

H.1 Identity as a Sheaf

Let X denote the space of contexts, which may include jurisdictions, physical locations, or social settings. An identity structure over this space can be modeled by a sheaf \mathcal{I} assigning to each open set $U \subseteq X$ the set $\mathcal{I}(U)$ of all identities valid within that context. The associated restriction maps express how identity information must behave when moving from a larger context to a smaller one, enforcing compatibility with the legal, physical, or social constraints that govern identity across overlapping regions. In this way, the sheaf formalism captures the locality, contextual dependence, and coherence requirements inherent in real-world identity systems.

Compatibility requires:

$$s_U|_V = s_V \quad \forall V \subset U.$$

H.2 Ledger Representation

Blockchain represents identity via keypairs:

$$\mathcal{K}(U) = \{\text{public keys active in } U\}.$$

Restrictions are trivial subset relations.

H.3 Sheaf Morphism Requirement

A morphism $\phi : \mathcal{I} \rightarrow \mathcal{K}$ between the identity sheaf and the keypair sheaf must satisfy several structural conditions. It must be compatible with all restriction maps so that identity information behaves consistently across overlapping jurisdictions or contexts. It must preserve equivalence classes of identities, ensuring that legally or socially identical persons are mapped to equivalent key-based representations. Furthermore, in contexts where identity is uniquely determined, the morphism must be injective on the corresponding stalks in order to avoid conflating distinct individuals.

H.4 No-Go Result

Theorem H.1. *There exists no sheaf morphism $\phi : \mathcal{I} \rightarrow \mathcal{K}$ that is simultaneously injective on stalks and compatible with real-world identity constraints.*

Reason. In practice, the mapping from legal or institutional identity to blockchain keypairs violates each of the required structural properties. Distinct legal identities may be represented by the same cryptographic key, while a single individual may control multiple independent keys without any canonical means of establishing equivalence. Jurisdiction-specific identity conditions have no analogue in the restriction maps of the ledger sheaf, since blockchain restrictions merely enforce syntactic continuity rather than legal or physical constraints. Moreover, ledger updates can occur entirely independently of real-world identity verification, allowing the on-chain state to diverge arbitrarily from the off-chain identity structure it is supposed to encode. Consequently, identity cannot be functorially embedded into the ledger structure, and no sheaf morphism with the required properties can exist.

Appendix I: Variational Derivation of the Useless-Work Criterion

This appendix derives the criterion for useless computation from a variational principle applied to entropy dynamics.

I.1 Action Functional for Entropy Reduction

Define an “entropy action”:

$$\mathcal{A}[S] = \int_{t_0}^{t_1} \int_{\Omega} (\partial_t S + \lambda \mathcal{C}(S)) w(x, t) d^3x dt,$$

where $\mathcal{C}(S)$ enforces physical constraints (continuity, conservation laws).

Systems that produce useful work must minimize \mathcal{A} .

I.2 Euler–Lagrange Equation

Taking the variation δS :

$$\delta \mathcal{A} = \int (\delta \partial_t S) w = [\delta S w]_{t_0}^{t_1} - \int \delta S \partial_t w.$$

If $\partial_t w = 0$ (weights fixed),

$$\frac{\delta \mathcal{A}}{\delta S} = 0 \Rightarrow \partial_t w = 0.$$

Thus minimizing action requires:

$$\partial_t S \leq 0.$$

I.3 Useless Computation Criterion

A computation (or infrastructure) satisfies:

$$\mathcal{V} = - \int \partial_t S w = 0 \iff \partial_t S = 0$$

and is useless if:

$$\mathcal{V} < 0 \quad (\text{entropy increasing}).$$

PoW, PoS satisfy $\partial_t S_{\text{real}} = 0$ but $\partial_t S_{\text{heat}} > 0$.

Thus:

$$\mathcal{V}_{\text{blockchain}} \approx 0, \quad \text{effective usefulness} < 0 \text{ when opportunity cost included.}$$

This formally justifies the core definition of “useless work.”

Appendix J: Token Economies and the Failure of Market Clearing

This appendix provides a formal microeconomic analysis showing why cryptocurrency markets rarely achieve efficient equilibrium, even under idealized assumptions.

J.1 Setup: Exchange Economy with Synthetic Assets

Consider an exchange economy with agents $i = 1, \dots, n$ who trade a cryptocurrency token T and a vector of real goods G .

Utility functions:

$$U_i(T, G) = u_i(G) + \gamma_i \cdot f(T).$$

The key feature is that $f(T)$ depends only on token quantity, not on real resource allocation; i.e. tokens have no direct productive role.

J.2 No Walrasian Equilibrium

Let p_T be token price and p_G a vector of real-good prices.

A Walrasian equilibrium requires excess demand $Z(T, G)$ satisfy $Z = 0$. But token demand satisfies:

$$Z_T = \sum_i \gamma_i f'(T_i),$$

which has no dependence on real-good allocation.

Thus, unless $f'(T_i) = 0$ for all agents (degenerate), token demand is unanchored.

Theorem J.1. In any token-based economy where token utility is independent of real goods, there exists no nontrivial Walrasian equilibrium.

This mirrors classical results that speculative or intrinsically valueless assets generate indeterminacy and non-clearing markets.

J.3 Instability Through Second-Order Price Feedback

Price dynamics follow:

$$\dot{p}_T = \alpha \cdot Z_T,$$

but Z_T is itself a function of future expected price profiles:

$$Z_T = \mathbb{E}[\Delta p_T].$$

Thus the system reduces to:

$$\dot{p}_T = \alpha \cdot \mathbb{E}[\dot{p}_T],$$

a self-referential, unstable recursion.

Cryptocurrency markets therefore exhibit persistent volatility, recurrent bubble formation, an inherent inability to converge toward stable equilibria, and a systematic failure to allocate real resources efficiently. Taken together, these properties formalize why token-based economies cannot function as effective market mechanisms and instead operate as persistent generators of speculative bubbles.

Appendix K: Epistemic Failure via Bayesian Network Analysis

This appendix gives a structural epistemic argument using Bayesian networks.

K.1 World-State Bayesian Network

Let $\mathcal{B}_{\text{world}}$ be a Bayesian network encoding real causal dependencies among variables X_1, \dots, X_n .

Each variable evolves under conditional distributions:

$$P(X_i | \text{Pa}(X_i)).$$

K.2 Ledger-State Bayesian Network

Ledger states form a trivial DAG:

$$\ell_{t+1} \leftarrow \ell_t,$$

with no causal arrows from the world.

Let $\mathcal{B}_{\text{chain}}$ encode:

$$P(\ell_{t+1} | \ell_t, \text{transactions}_t).$$

K.3 No Information Flow Theorem

Theorem K.1 (Epistemic Isolation). If the ledger contains no causal parents of any real-world variable, then no Bayesian updating based on ledger states can reduce uncertainty about the real world.

Proof. Under the Markov condition, each real-world variable X_i is independent of the ledger state ℓ_t , since the parent set $\text{Pa}(X_i)$ contains no ledger nodes. Formally,

$$X_i \perp\!\!\!\perp \ell_t.$$

It follows that posterior updates conditioned on the entire history of the ledger do not alter the distribution of X_i :

$$P(X_i | \ell_0, \dots, \ell_t) = P(X_i).$$

Thus the ledger provides no epistemic evidence about any real-world variable, and no Bayesian inference over ledger data can reduce real-world uncertainty.

K.4 Consequence: Zero Epistemic Value

Because blockchain ledgers encode no causal structure of the environment, they cannot function as epistemically meaningful infrastructures for identity, provenance, scientific data, or any knowledge system that depends on causal constraints. Their observable states remain decoupled from the physical and institutional processes they purport to represent, rendering them informationally sterile with respect to real-world inference.

Appendix L: Complexity of Semantic Verification

We now demonstrate that verifying the semantic correctness of ledger states, when measured against real-world constraints, is NP-hard and in many cases even more computationally demanding.

L.1 Semantic Constraint Satisfaction

Real-world asset consistency, identity correctness, and other semantic requirements can be expressed as a constraint satisfaction problem (CSP)

$$\mathcal{C} = \{c_1, \dots, c_m\},$$

defined over variables that describe world events or states. A ledger state ℓ can be interpreted as a proposed assignment to these variables. Semantic verification then asks whether the ledger state satisfies all constraints in \mathcal{C} :

$$\text{Is } \ell \models \mathcal{C}?$$

In other words, the question is whether the syntactic ledger state aligns with the semantic conditions imposed by the legal, physical, or institutional reality it is intended to encode. As we show in the subsequent sections, determining this alignment is computationally intractable in general.

L.2 NP-Hardness Result

Because general constraint satisfaction problems (CSPs) are NP-complete, any verification task that requires solving an arbitrary CSP inherits this computational hardness. In the context of blockchain systems, the ledger's syntactic structure contains no representation of real-world semantic constraints. As a result, verifying whether a given ledger state is semantically correct requires checking consistency against off-chain information, which includes legal, physical, and institutional facts that the ledger does not and cannot encode.

Theorem L.1 (Semantic Verification Hardness). Verifying that a blockchain ledger correctly represents a world-state satisfying all relevant constraints is NP-hard.

Sketch of proof. Consider an arbitrary CSP instance expressed as a collection of constraints on world-state variables. This instance can be mapped directly into a requirement that a corresponding ledger state encode a world consistent with those constraints. Since the ledger contains only syntactic data and lacks any internal structure for enforcing or even representing semantic relations, the task of determining whether the ledger state is semantically valid becomes equivalent to checking whether the CSP instance is satisfiable. Thus the verification problem is at least as hard as the original CSP. Consequently, blockchain technologies do not simplify the problem of semantic verification; they merely re-express the original intractable decision problem in a setting where none of the off-chain constraints are natively enforced.

L.3 Impossibility of Polynomial-Time Auditing

Since real-world audits must check semantic correctness, and ledger structure adds no constraints, the system provides no computational advantage.

This formalizes that blockchain cannot reduce verification costs for complex assets or identities.

Appendix M: Stochastic Dynamics of Bubble Formation

This appendix derives the stochastic differential equation (SDE) structure responsible for persistent bubbles in token markets.

M.1 Price as Drift Plus Speculative Feedback

Let token price follow:

$$dp_t = \mu_t dt + \sigma dW_t,$$

but μ_t is endogenous:

$$\mu_t = \alpha \mathbb{E}_t[\Delta p_t].$$

Thus:

$$dp_t = \alpha \mathbb{E}_t[\Delta p_t] dt + \sigma dW_t.$$

M.2 Infinite-Reactivity Feedback Loop

Let beliefs evolve by:

$$\mathbb{E}_t[\Delta p_t] = k(p_t - \bar{p}),$$

for reference price \bar{p} .

Then:

$$dp_t = \alpha k(p_t - \bar{p})dt + \sigma dW_t,$$

an Ornstein–Uhlenbeck process with *wrong sign* (unstable drift).

M.3 Condition for Bubble Explosion

Solution:

$$p_t = \bar{p} + (p_0 - \bar{p})e^{\alpha kt} + \sigma \int_0^t e^{\alpha k(t-s)} dW_s.$$

If $\alpha k > 0$, price diverges exponentially. Token markets generically satisfy this because expectation is linked to recent price action.

Thus, bubbles in token markets are not incidental anomalies but structural features of the underlying stochastic dynamics. The instability arises from the positive-feedback relationship between price expectations and subsequent price movements, which ensures that speculative corrections do not restore a long-run equilibrium and that volatility remains an endemic property of the system rather than a transient deviation.

M.4 Blockchain Cannot Suppress Bubbles

Blockchain infrastructures lack any mechanism capable of counteracting this form of stochastic instability. Because they impose no upper bounds on price excursions, provide no semantic grounding that could anchor prices to real-world value, and fail to supply any fundamental valuation metric, the unstable Ornstein–Uhlenbeck process described above cannot be stabilized within a blockchain-based economy. The architecture is therefore intrinsically incapable of suppressing bubble formation or preventing runaway speculative dynamics.

Blockchain therefore functions as a bubble generator by design.

Appendix N: Semantic Information Geometry and Useful Work

Here we formalize the curvature structure underlying the notion of "useful work."

N.1 Semantic Manifold and Fisher Metric

Let \mathcal{M} be a manifold of semantic states with coordinates θ .

Define Fisher metric:

$$g_{ij}(\theta) = \mathbb{E} [\partial_i \log p(x|\theta) \partial_j \log p(x|\theta)].$$

Curvature R measures how rapidly predictions change in response to perturbations.

N.2 Entropy-Reducing Flows

Useful work corresponds to flows on \mathcal{M} that reduce expected uncertainty:

$$\dot{\theta}^i = -g^{ij}\partial_j \mathcal{S}$$

(steepest descent on entropy landscape).

N.3 Blockchain as Zero-Curvature Semantic System

Ledger states depend only on syntax. Thus:

$$p(x|\theta) = p(x) \quad \forall \theta,$$

so Fisher information is zero:

$$g_{ij} = 0, \quad R = 0.$$

Thus no gradient flow exists:

$$\dot{\theta} = 0.$$

This proves that blockchains cannot induce movement toward lower-entropy semantic configurations.

N.4 Interpretation

Semantic information geometry makes clear why blockchains are incapable of updating or maintaining coherent semantic world-models. The space of ledger states forms a flat manifold with no intrinsic curvature, and the update rules defined by the consensus protocol move

exclusively along syntactic directions that preserve this flatness. Because no geometric pathway connects these syntactic trajectories to the curved manifolds that represent genuine semantic structure, the ledger cannot participate in the creation, refinement, or maintenance of meaningful information. The information-geometric analysis therefore completes the derivation of the conclusion that blockchain systems perform useless work: they update surface-level representations without any mechanism for generating or transforming semantic structure.

References

- Bamieh, Bassam and Dennice F. Gayme (2012). “The Price of Synchrony: Resistive Losses due to Phase Synchronization in Power Networks”. In: *IEEE Transactions on Automatic Control* 57.9, pp. 2281–2296.
- Bennett, Charles H. (1973). “Logical Reversibility of Computation”. In: *IBM Journal of Research and Development* 17.6, pp. 525–532.
- Cheng, Newman, Gordon Broadbent, and William Chappell (2025). “Cognitive Loop via In-Situ Optimization: Self-Adaptive Reasoning for Science”. In: arXiv: 2508.02789 [cs.AI].
- Easley, David, Maureen O’Hara, and Soumya Basu (2019). “From Mining to Markets: The Evolution of Bitcoin Transaction Fees”. In: *Management Science* 65.12, pp. 5593–5608.
- England, Jeremy L. (2013). “Statistical Physics of Self-Replication”. In: *The Journal of Chemical Physics* 139.12, p. 121923.
- Friston, Karl (2005). “A Theory of Cortical Responses”. In: *Philosophical Transactions of the Royal Society B* 360.1456, pp. 815–836.
- Gencer, Adem Efe et al. (2018). “Decentralization in Bitcoin and Ethereum: A Measurement Study”. In: *2018 IEEE Symposium on Security and Privacy (SP)*, pp. 683–698.
- Ghrist, Robert (2014). *Elementary Applied Topology*. Createspace.
- Grossberg, Stephen (1976). “Adaptive Pattern Classification and Universal Recoding: I. Parallel Development and Coding of Neural Feature Detectors”. In: *Biological Cybernetics* 23.3, pp. 121–134.
- (1987). “Competitive Learning: From Interactive Activation Models to Adaptive Resonance Theory”. In: *Cognitive Science* 11.1, pp. 23–63.
- (1990). “Neural Networks and Natural Intelligence: Models of Adaptive Resonance and Self-Organization”. In: *Proceedings of the IEEE* 78.9, pp. 1464–1481.
- Hansen, Jakob and Robert Ghrist (2018). “Toward a Spectral Theory of Cellular Sheaves”. In: arXiv: 1805.11973 [math.AT].
- Landauer, Rolf (1961). “Irreversibility and Heat Generation in the Computing Process”. In: *IBM Journal of Research and Development* 5.3, pp. 183–191.
- Lynch, Nancy A. (1996). *Distributed Algorithms*. Morgan Kaufmann.

- Millidge, Beren, Anil K. Seth, and Christopher L. Buckley (2021). “Predictive Coding: A Theoretical and Experimental Review”. In: *Neural Computation* 33.11, pp. 2450–2502.
- Narayanan, Arvind et al. (2016). *Bitcoin and Cryptocurrency Technologies*. Princeton University Press.
- Olfati-Saber, Reza, J. Alex Fax, and Richard M. Murray (2007). “Consensus and Cooperation in Networked Multi-Agent Systems”. In: *Proceedings of the IEEE* 95.1, pp. 215–233.
- Schrödinger, Erwin (1944). *What Is Life?* Cambridge University Press.
- Seely, Jeffrey (2025). “Sheaf Cohomology of Linear Predictive Coding Networks”. In: arXiv: 2511.11092 [cs.LG].
- Sohl-Dickstein, Jascha, Eric Weiss, Niru Maheswaranathan, and Surya Ganguli (2015). “Deep Unsupervised Learning Using Nonequilibrium Thermodynamics”. In: *International Conference on Machine Learning*, pp. 2256–2265.
- Song, Yang and Stefano Ermon (2019). “Generative Modeling by Estimating Gradients of the Data Distribution”. In: *Advances in Neural Information Processing Systems* 32.
- Spivak, David I. (2014). *Category Theory for the Sciences*. MIT Press.
- Tong, Shangyuan, Nanye Ma, Saining Xie, and Tommi Jaakkola (2025). “Flow Map Distillation Without Data”. In: arXiv: 2511.19428 [cs.LG].
- Touchette, Hugo and Seth Lloyd (2000). “Information-Theoretic Limits of Control”. In: *Physical Review Letters* 84.6, pp. 1156–1159.